

Deep Learning for Tabular Data: An Exploratory Study

by

Jan André Marais



*Thesis presented in partial fulfilment of the requirements for
the degree of Master of Commerce (Mathematical Statistics)
in the Faculty of Economic and Management Sciences at
Stellenbosch University*

Supervisor: Dr. S. Bierman

December 2018

The financial assistance of the National Research Foundation (NRF) towards this research is hereby acknowledged. Opinions expressed and conclusions arrived at, are those of the author and are not necessarily to be attributed to the NRF.

Declaration

By submitting this thesis electronically, I declare that the entirety of the work contained therein is my own, original work, that I am the sole author thereof (save to the extent explicitly otherwise stated), that reproduction and publication thereof by Stellenbosch University will not infringe any third party rights and that I have not previously in its entirety or in part submitted it for obtaining any qualification.

Date:

Copyright © 2018 Stellenbosch University
All rights reserved.

Abstract

Deep Learning for Tabular Data: An Exploratory Study

J. A. Marais

Thesis: MCom (Mathematical Statistics)

December 2018

English abstract.

Uittreksel

Diepleer Tegnieke vir Gestruktrueerde Data

(“Deep Learning for Tabular Data: An Exploratory Study”)

J. A. Marais

Tesis: MCom (Wiskundige Statistiek)

Desember 2018

Afrikaans abstract

Acknowledgements

I would like to express my sincere gratitude to the following people and organisations:

- Dr. S. Bierman for her guidance and patience as a supervisor and for allowing me freedom in the choice of research directions.
- The National Research Foundation (NRF) for financial support under grant number: UID xxx. *Opinions expressed and conclusions arrived at, are those of the author and are not necessarily to be attributed to the NRF.*
- The UCI Machine Learning Repository (Dheeru and Karra Taniskidou, 2017) for hosting a platform to share datasets.
- My parents and close family for believing in me and motivating me.
- Most importantly, my partner, for her never-ending support, love and understanding.

Contents

Declaration	i
Abstract	ii
Uittreksel	iii
Acknowledgements	iv
Contents	v
List of Figures	ix
List of Tables	xii
List of Abbreviations and/or Acronyms	xiii
Nomenclature	xiv
1 Introduction	1
1.1 Deep Learning	1
1.2 Tabular Data	3
1.3 Challenges of Deep Learning for Tabular Data	5
1.4 Overview of Statistical Learning Theory	7
1.5 Outline	12
2 Neural Networks	14
2.1 Introduction	14
2.2 The Structure of a Neural Network	15
2.2.1 Neurons and Layers	15
2.2.2 Activation Functions	18

2.2.3	Size of the Network	20
2.3	Training a Neural Network	21
2.3.1	Optimisation	21
2.3.2	Optimisation Example	23
2.3.3	Backpropogation	24
2.3.4	Weight initialisation	27
2.3.5	Basic Regularisation	28
2.3.6	Adaptive Learning Rates and Annealing	28
2.4	Representation Learning	29
3	Deep Learning	34
3.1	Introduction	34
3.2	Denoising Autoencoders	35
3.3	Pretraining and Transfer Learning	37
3.4	More Regularisation	39
3.4.1	Dropout	39
3.4.2	Data Augmentation	39
3.5	Modern Architectures	41
3.5.1	Normalisation	41
3.5.2	Skip-connections	42
3.5.3	Embeddings	43
3.5.4	Attention	45
3.6	Super-Convergence	46
3.7	Model Interpretation	50
3.7.1	Neural Network Specific	50
3.7.2	Model Agnostic	51
4	Deep Learning for Tabular Data	54
4.1	Introduction	54
4.2	Input Representation	55
4.2.1	Numerical Features	55
4.2.2	Categorical Features	57
4.2.3	Combining Features	61
4.3	Learning Feature Interactions	61
4.3.1	Attention	64
4.3.2	Self-Normalising Neural Networks	65
4.4	Sample Efficiency	66

4.4.1	Data Augmentation	66
4.4.2	Unsupervised Pretraining	69
4.4.3	Regularisation	70
4.5	Interpretation	71
4.6	Hyperparameter Selection	75
5	Experiments	78
5.1	Introduction	78
5.2	Datasets	79
5.3	General Methodology	80
5.3.1	Loss Function and Evaluation Metric	80
5.3.2	Cross-validation	80
5.3.3	Preprocessing	82
5.3.4	General Hyperparameters	82
5.4	Input Representation	83
5.4.1	Embeddings Sizes	83
5.5	Feature Interactions	84
5.5.1	Attention	84
5.5.2	SeLU	85
5.5.3	Skip-Connections	85
5.6	Sample Efficiency	86
5.6.1	Data Augmentation	87
5.6.2	Unsupervised Pretraining	88
5.7	General Thoughts	90
6	Conclusion	91
6.1	Summary	91
6.2	Limitations	93
6.3	Promising Future Directions	94
Appendices		96
A	Datasets	97
A.1	Adult Dataset	97
B	Hyperparameter Search	101
B.1	Width and Depth of Network	101
B.2	Dropout	101

CONTENTS	viii
C Software and Code	103
C.1 Development Environment	103
C.2 Code and Reproducibility	103
Bibliography	104

List of Figures

1.1	The exponential growth of published papers and Google search terms containing the term <i>Deep Learning</i> . Sources: Google Trends, Semantic Scholar	2
1.2	Linear model on simple binary classification dataset.	11
2.1	Comparison of a biological (a) and an artificial (b) neuron.	16
2.2	A simple neural network accepting p -sized inputs, with one hidden layer which has two neurons.	17
2.3	Plots of various activation functions (a) and their local derivatives (b).	18
	(a)	18
	(b)	18
2.4	Plots of the gradient descent example. (a) The data points in input space. The shades in the background represent the class division in input space, with the decision boundary determined by linear least squares estimation. The dashed lines represent the decision boundaries learned at different iterations. (b) The loss calculated at each iteration.	24
2.5	Simple dataset with two linearly inseparable classes.	31
2.6	Decision boundary of 1-layer neural network.	31
2.7	Decision boundary of 2-layer neural network.	32
2.8	Hidden representation of 2-layer neural network.	32
3.1	A simple single hidden layer autoencoder with 4-dimensional inputs and 2 neurons in the hidden layer.	36
3.2	Visualising the first layer convolutional filters learned by a neural network in a large image dataset.	38
3.3	The effect dropout has on connections between neurons.	40

3.4	Example data augmentations for images.	40
3.5	Diagram conceptualising a skip-connection.	43
3.6	Learned word embeddings in a 2-dimensional space.	44
3.7	Attention applied to image captioning.	45
3.8	Attention applied to machine translation.	46
3.9	Reduced training iterations and improved performance with super-convergence principle.	48
3.10	The learning rate schedule of the 1Cycle policy.	48
3.11	An example output of a learning rate range test.	49
4.1	Effect of normalisation on continuous variables.	56
(a)	Original	56
(b)	Gaussian Norm	56
(c)	Power Norm	56
4.2	PCA of the Education entity embedding weight matrix.	60
4.3	Combined representation of continuous and categorical features.	62
4.4	Illustration of points created by mixup augmentation.	68
4.5	The attention weights visualised for a single observation in the dataset.	72
4.6	Permutation importance plot of a neural network trained on the Adult dataset.	73
4.7	Feature importance obtained from a boosted trees model trained on the neural network predictions.	74
4.8	Constant learning rate vs the 1Cycle schedule.	76
4.9	A learning rate range test with different weight decays.	77
4.10	A full training run with different weight decays.	77
5.1	5-Fold Cross-validation dataset split schematic.	81
5.2	Effect of the embedding size if all categorical features are mapped to the same number of dimensions.	83
5.3	Effect of variable sizes on the performance of the model.	84
5.4	Comparing the attention mechanism with a simple MLP.	85
5.5	The average performance of ReLU and SeLU activation functions for shallow and deep networks over every epoch.	86
5.6	The average performance of ReLU and SeLU activation functions for shallow and deep networks.	86
5.7	Average performance at each epoch for shallow and deep neural networks, with and without skip connections.	87

5.8	Overall performance of the skip-connections used in a shallow and deep neural network.	87
5.9	Effect of the number of training samples on the performance of neural networks.	88
5.10	Average performance of models with various mixup and weight decay parameters.	88
5.11	Performance per epoch for models with different weight decays and mixup ratios.	89
5.12	The effect of pretraining on the classifier's performance.	90
A.1	Histograms for each of the continuous features in the Adult dataset.	99
A.2	Bar plot for each of the categorical features in the Adult dataset. .	100
B.1	Effect of the layer width and network depth on the performance on the Adult dataset.	102
B.2	The effect of dropout on wide and narrow neural networks.	102

List of Tables

1.1	Preview of the Adult dataset.	4
4.1	Swap Noise Example.	67

List of Abbreviations and/or Acronyms

ANN	Artificial Neural Network
CNN	Convolutional Neural Network
CTR	Click-through Rate
CV	Computer Vision
DL	Deep Learning
EHR	Electronic Health Records
GAN	Generative Adversarial Network
kNN	k -Nearest Neighbour
mAP	Mean Average Precision
ML	Machine Learning
MLP	Multi-layer Perceptron
NLP	Natural Language Processing
NN	Neural Network
RNN	Recurrent Neural Network
SGD	Stochastic Gradient Descent
SotA	State-of-the-Art
VAE	Variational Autoencoder

Nomenclature

N	number of observations in a dataset
p	input dimension or the number of features for an observation
K	number of labels in a dataset
\mathbf{x}	p -dimensional input vector $(x_1, x_2, \dots, x_p)^\top$
λ	label
\mathcal{L}	complete set of labels in a dataset $\mathcal{L} = \{\lambda_1, \lambda_2, \dots, \lambda_K\}$
Y	labelset associated with \mathbf{x} , $Y \subseteq \mathcal{L}$
\hat{Y}	predicted labelset associated with \mathbf{x} , $\hat{Y} \subseteq \mathcal{L}$, produced by $h(\cdot)$
\mathbf{y}	K -dimensional label indicator vector, $(y_1, y_2, \dots, y_K)^\top$, associated with observation \mathbf{x}
$(\mathbf{x}_i, Y_i)_{i=1}^N$	multi-label dataset with N observations
D	dataset
$h(\cdot)$	multi-label classifier $h : \mathbb{R}^p \rightarrow 2^{\mathcal{L}}$, where $h(\mathbf{x})$ returns the set of labels for \mathbf{x}
θ	set of parameters for $h(\cdot)$
$\hat{\theta}$	set of parameters for $h(\cdot)$ that optimise the loss function
$L(\cdot, \cdot)$	loss function between predicted and true labels
$f(\cdot)$	label prediction module, $f : \mathbb{R}^p \rightarrow \mathbb{R}^K$
$t(\cdot)$	thresholding function, $t : \mathbb{R}^K \rightarrow \{0, 1\}^K$
$\mathcal{N}(\mathbf{x})$	points in the input space neighbourhood of \mathbf{x}

Chapter 1

Introduction

1.1 Deep Learning

This thesis is concerned with the study of *deep learning* approaches to solve *machine learning* (ML) tasks. More specifically, our interest lies in machine learning tasks that may be solved using tabular data inputs. The deep learning field is an extension of the class of machine learning algorithms called *Artifical Neural Networks* (NNs). Whereas until relatively recently, the neural network field was not a very active research field, rapid development in computing power and the growing abundance of data lead to advances in neural network optimisation and architecture. These advances constitutes the deep learning field as we know it today (Lecun *et al.*, 2015).

Currently, deep learning is receiving a remarkable amount of attention, both in research and in practice (see Figure 1.1). Much of the deep learning hype stems from the tremendous value neural networks have shown in application areas such as *computer vision* (Hu *et al.*, 2017), audio processing (Battenberg *et al.*, 2017), and *natural language processing* (NLP) (Devlin *et al.*, 2018). In these application areas, deep learning methods have reached a maturity level sufficient to be able to run these systems in a production or commercial environment. Examples of the application of deep learning in commercial applications include voice assistants like Amazon Alexa (Sarikaya, 2017), face recognition with Apple iPhones¹, and language translation with Google (Wu *et al.*, 2016).

¹https://www.apple.com/business/site/docs/FaceID_Security_Guide.pdf

²<https://trends.google.com/trends/explore?date=all&q=deep%20learning>

³<https://www.semanticscholar.org/search?year%5B0%5D=2000&year%5B1%5D=>




Figure 1.1: The exponential growth of published papers and Google search terms containing the term *Deep Learning*. Sources: Google Trends², Semantic Scholar³

One of the most attractive attributes of deep learning is its ability to model almost any input-output relationship. This has lead to the use of deep learning in a very wide array of applications.

For example, deep learning has been used to generate art (Gatys *et al.*, 2015) and music (Mogren, 2016), to control various modules in autonomous cars (Fridman *et al.*, 2017), to play video games (Mnih *et al.*, 2013), to recommend movies (Covington *et al.*, 2016), to improve the quality of images (Shi *et al.*, 2016), and to beat the world's best Go player (Silver *et al.*, 2017).

A common characteristic of all of the above deep learning applications, is that the data used to construct them, contain the same type of values or measurements. That is, in computer vision the data are pixel values, in NLP the data are words and in audio processing the data represent sound waves. This is not a criterion for deep learning algorithms to be successful, but may be viewed as a driver for their sucess in these application domains. It is simpler to model data consisting of the same type of measurements, since each input feature can be treated the same. Furthermore in the above deep learning

²2019&q=%22deep%20learning%22&sort=relevance

applications, it is found that in each of these domains, universal patterns exist. This allows for knowledge to be transferred between tasks belonging to the same domain. The knowledge to be transferred is both knowledge aquired by humans, and that learned by a deep learning model. For example, in computer vision, advancements in classifying pictures of pets will most likely also be facilitate improved identification of tumors in X-rays. That is, patterns learned by a deep learning model when attempting one task, may also be useful in a different, but related task (see *Transfer Learning*).

A data domain in which deep learning does not flourish is that of tabular data. A *tabular dataset* can be represented by a two-dimensional table, where each of the rows of the table corresponds to one observation and where each column denotes an individual meaningful feature. We further explain the use of tabular data in Section 1.2 below.

Some research have recently been done on the use of deep learning models for tabular data (Shavitt and Segal, 2018, Song *et al.* (2018)). However, state-of-the-Art (SotA) results are reported only rarely (de Brébisson *et al.*, 2015) (and in this competition⁴). Therefore it can be said that the area is nowhere near as mature or receiving as much attention as is the case with deep learning for computer vision or for NLP. In a comprehensive study in the paper by (Fernández-Delgado *et al.*, 2014), it was found that ML tasks that make use of tabular data are typically more effectively solved using tree-based methods. This is also evident when one considers the winning solutions of relevant Kaggle competitions⁵. A possible explanation for the superior performance of tree-based metods, is the heterogeneity of tabular data (Shavitt and Segal, 2018), which forms part of the discussion in the next section.

1.2 Tabular Data

In this section, we make use of the so-called Adult⁶ dataset in order to discuss the use of tabular data. The reader may refer to Table 1.1 for an extract of this dataset. The data were collected during an American census with the aim of predicting whether or not an individual earns more than \$50,000 a year.

⁴<https://www.kaggle.com/c/porto-seguro-safe-driver-prediction/discussion/44629>

⁵<https://www.kaggle.com>

⁶<http://archive.ics.uci.edu/ml/datasets/Adult>

	age	occupation	education	race	sex	>=50k
1	49		Assoc-acdm	White	Female	1
2	44	Exec-managerial	Masters	White	Male	1
3	38		HS-grad	Black	Female	0
4	38	Prof-specialty	Prof-school	Asian-Pac-Islander	Male	1
5	42	Other-service	7th-8th	Black	Female	0
6	20	Handlers-cleaners	HS-grad	White	Male	0

Table 1.1: Preview of the Adult dataset.

Table 1.1 represents a typical tabular dataset, where the columns contain measurements on different features. Therefore each column may contain different data types. Some columns may consist of continuous measurements, whereas other columns contain discrete or categorical measurements. Hence the columns are heterogeneous with respect to data types. Furthermore, in tabular data, the rows and columns occur in no particular order. This of course stands in contrast to image or text data.

Many important ML applications make use of tabular data. Some of these applications are listed below:

- Various tasks that make use of Electronic Health Records. These include the prediction of in-hospital mortality rates, and prolonged length of stay (Rajkomar *et al.*, 2018).
- Recommender systems for items like videos (Covington *et al.*, 2016) or property listings (Halder *et al.*, 2018).
- Click-through rate (CTR) prediction in web applications, *i.e.* predicting which item a user will click on next (Song *et al.*, 2018).
- Predicting which clients are at risk of defaulting on their accounts⁷.
- Predicting store sales (Guo and Berkhahn, 2016).
- Drug discovery (Klambauer *et al.*, 2017).

Tabular datasets come in all shapes and sizes. The number of rows can range from hundreds to millions, and the number of columns also has no limits. It is not unusual for tabular datasets to be noisy. A proportion of the observations may have missing features and/or incorrect values. Continuous measurements may be based upon vastly different scales, some containing outliers, whereas categorical features may have high cardinality which leads to sparse data.

⁷<https://www.kaggle.com/c/loan-default-prediction>

During the construction of models for tabular datasets, the most important step in terms of leading to improvements in the model performance, is pre-processing and manipulation of the input features (Rajkomar *et al.*, 2018). This includes data merging, customising, filtering and cleaning. In a process called feature engineering one strives to create new features from the original features based on some domain knowledge. The idea is that such engineered features enables a model to learn interactions between features, thereby facilitating more accurate prediction. Feature engineering is an extremely laborious process with no clear recipe to follow and therefore typically cannot successfully be implemented without some domain expertise.

Ensemble methods based upon trees are currently viewed as the most effective machine learning models in the face of tabular datasets. As mentioned above, a possible reason for this may be their robustness to different feature scales and data types, linked with their ability to effectively model interactions among features having different data types.

Indeed, in the context of tabular data, classical neural network approaches are no match for tree ensembles. Although the deep learning field has advanced and matured a lot in recent years, it is not yet clear how to leverage these modern techniques to effectively build and train deep neural networks on tabular datasets. In this thesis we explore ways of doing so. By reviewing the most recent literature on the topic, and through empirical work, we aim to summarise best practices when using deep learning for tabular data.

1.3 Challenges of Deep Learning for Tabular Data

Some of the challenges of deep learning for tabular data have been alluded to in earlier sections of this chapter. These will form the framework for the literature review which follows later on, therefore they are summarised below, in terms of relevant questions to ask when applying deep learning in the context of tabular data:

- **How should input features be represented numerically?:** We have mentioned that tabular data consist of mixed feature types, *i.e.* a combination of categorical and continuous features. The question here

relates to how these features should be processed and presented to the model during training.

- **How can we exploit feature interactions?:** Once we have found the optimal feature representation for all feature types, we will need a way to effectively learn the interactions between them and also a way to learn how they relate to the target. This is a crucial step towards the effective application of deep learning models to tabular data.
- **How can we be more sample efficient?:** Tabular datasets are typically smaller than the datasets used in computer vision and in NLP. Moreover, no general large dataset with universal properties exists to be used by a model to learn from (as is the case in for example image classification). Thus a key challenge is to facilitate learning from less data.
- **How do we interpret model decisions?:** The use of deep learning is often restricted by its perceived lack of interpretability. Therefore we need ways of explaining the model output in order for it to be useful in many applications.

Clearly there are plenty of considerations when it comes to using deep learning for tabular data. The main objective of this thesis is to find the best ways of overcoming the above challenges. Towards this objective, the study should lead to a thorough understanding of the *status quo* of the field, and of the necessary factors in order to ensure deep learning to be as effective in other data domains as it currently is in computer vision and NLP.

The study is divided into two parts. We start by first providing an overview of the relevant literature. Subsequently, we make use of experimental work in order to compare various deep learning algorithms (and possible improvements) on relevant datasets. Here an important aim will be to ensure our experiments to be *rigorous*. The importance of rigorous research has relatively recently again been emphasised during an NIPS talk⁸, during which researchers in the deep learning field have been criticised for the growing gap between the understanding of its techniques and practical successes. Currently much more emphasis is placed on the latter. The speakers urged the deep learning community to be more rigorous in their experiments where, for them, the most important part of rigor is better empiricism, not more mathematical theories. Better

⁸Talk given at NIPS2017 - <https://www.youtube.com/watch?v=Qi1Yry33TQE>

empiricism in classification may include, for example, practices such as using cross-validation to estimate the generalisation ability of a model, and reporting standard errors. Empirical studies should involve more than simply attempting to beat the benchmark, for example where possible, they should also involve simple experiments that facilitate understanding why some algorithms are successful, while others are not.

In addition, we want the empirical work in this study to be as reproducible as possible. This aspect is often overlooked. However it is a crucial aspect, ensuring transparent and accountable reporting of results. Reproducibility adds to the value of research, since without it, researchers are not able to build on each other's work. There all code, data and necessary documentation in order to reproduce the experiments that were done in this thesis will be available⁹.

Having stated the objectives of this study, we now turn to a discussion of the fundamental concepts of Statistical Learning Theory. This is followed by a more detailed summary of the thesis.

1.4 Overview of Statistical Learning Theory

Machine- or statistical learning algorithms (used interchangeably) are used to perform certain tasks that are too difficult to solve with fixed rule-based programs. Hence statistical learning algorithms are able to use data in order to learn how to perform difficult tasks. For an algorithm to learn from data means that it can improve its ability to perform an assigned *task* with respect to some *performance measure*, by processing *data*. In this section we discuss some of the important types of tasks, data and performance measures in the statistical learning field.

A learning task describes the way in which an algorithm should process an observation. An observation is a collection of features that have been measured, corresponding to some object or event that we want the system to process, for example an image. We will represent an observation by a vector $\mathbf{x} \in \mathbb{R}^p$, where each element x_j of the vector is an observed value of the j -th feature, $j = 1, \dots, p$. For example, the features of an image are usually the color intensity values of the pixels in the image.

Many kinds of tasks can be solved using statistical learning. One of the most common learning tasks is that of *classification*, where it is expected of an

⁹Shared publicly at <https://github.com/jandremarais/tabularLearner>

algorithm to determine which of K categories an input belongs to. In order to complete the classification task, the learning algorithm is usually asked to produce a function $f : \mathbb{R}^p \rightarrow \{1, \dots, K\}$. When $y = f(\mathbf{x})$, the model assigns an input described by the vector \mathbf{x} to a category identified by the numeric code y , called the *output* or *response*. In other variants of the classification task, f may output a probability distribution over the possible classes.

Regression is another main learning task and requires the algorithm to predict a continuous value given some input. This task requires a function $f : \mathbb{R}^p \rightarrow \mathbb{R}$, where the only difference between regression and classification is the format of the output.

Learning algorithms learn such tasks by observing a relevant set of data points, *i.e.* a dataset. A dataset containing N observations of p features is commonly denoted by a data matrix $X : N \times p$, where each row represents a different observation and where each column corresponds to a different feature of the observations, *i.e.*

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N1} & x_{N2} & \dots & x_{Np} \end{bmatrix}.$$

Often the dataset includes annotations for each observation in the form of a label (*i.e.* in classification) or in the form of a target value (*i.e.* in regression). These N annotations are represented by the vector \mathbf{y} , where the element y_i is associated with the i -th row of X . Therefore the response vector may be denoted by

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}.$$

Note that in the case of multiple labels or targets, a matrix representation $Y : N \times K$ is required.

Statistical learning algorithms can be divided into two main categories, *viz.* *supervised* and *unsupervised* algorithms. This categorisation is determined by the presence (or absence) of annotations in the dataset to be analysed. Unsupervised learning algorithms learn from data consisting only of features,

X , and are used to find useful properties and structure in the dataset (see Hastie *et al.*, 2009, Ch. 14). On the other hand, supervised learning algorithms learn from datasets which consist of both features and annotations, (X, Y) , with the aim to model the relationship between them. Therefore, both classification and regression are considered to be supervised learning tasks.

In order to evaluate the ability of a learning algorithm to perform its assigned task, we have to design a quantitative performance measure. For example, in a classification task we are usually interested in the accuracy of the algorithm, *i.e.* the percentage of times that the algorithm assigns the correct classification. We are mostly interested in how well the learning algorithm performs on data that it has not seen before, since this demonstrates how well it will perform in real-world situations. Thus we typically evaluate the algorithm on a *test set* of data points. This dataset is independent of the *training set* of data points that used during the learning process.

For a more concrete example of supervised learning, and keeping in mind that the linear model is one of the main building blocks of neural networks, consider the learning task underlying *linear regression*. The objective here is to construct a system which takes a vector $\mathbf{x} \in \mathbb{R}^p$ as input and which predicts the value of a scalar $y \in \mathbb{R}$ as response. In the case of linear regression, we assume the output be a linear function of the input. Let \hat{y} be the predicted response. We define the output to be

$$\hat{y} = \hat{\mathbf{w}}^T \mathbf{x},$$

where $\hat{\mathbf{w}} = [w_0, w_1, \dots, w_p]$ denotes a vector of parameters and where $\mathbf{x} = [1, x_1, x_2, \dots, x_p]$. Note that an intercept is included in the model (also known as a *bias* in machine learning). The parameters are values that control the behaviour of the system. We can think of them as a set of *weights* that determine how each feature affects the prediction. Hence the learning task can be defined as predicting y from \mathbf{x} through $\hat{y} = \hat{\mathbf{w}}^T \mathbf{x}$.

We of course need to define a performance measure to evaluate the linear predictions. For a set of observations, an evaluation metric tells us how (dis)similar the predicted output is to the actual response values. A very common measure of performance in regression is the *mean squared error* (MSE), given by

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2.$$

The process of learning from the data (or fitting a model to the data) can be reduced to the following optimisation problem: find the set of weights, $\hat{\mathbf{w}}$, which produces a $\hat{\mathbf{y}}$ that minimises the MSE. Of course this problem has a closed form solution and can quite trivially be found by means of *ordinary least squares* (OLS) (see Hastie *et al.*, 2009, p. 12). However, we have mentioned that we are more interested in the algorithm's performance evaluated on a test set. Unfortunately the least squares solution does not guarantee the solution to be optimal in terms of the MSE on a test set, rendering statistical learning to be much more than a pure optimisation problem.

The ability of a model to perform well on previously unobserved inputs is referred to as its *generalisation* ability. To be able to fit a model that generalises well to new unseen data cases is the key challenge of statistical learning. One way of improving the generalisation ability of a linear regression model is to modify the optimisation criterion J , to include a *weight decay* (or *regularisation*) term. That is, we want to minimise

$$J(\mathbf{w}) = MSE_{\text{train}} + \lambda \mathbf{w}^T \mathbf{w},$$

where $J(\mathbf{w})$ now expresses preference for smaller weights. The parameter λ is non-negative and needs to be specified ahead of time. It controls the strength of the preference by determining how much influence the penalty term, $\mathbf{w}^T \mathbf{w}$, has on the optimisation criterion. If $\lambda = 0$, no preference is imposed, and the solution is equivalent to the OLS solution. Larger values of λ force the weights to decrease, and thus referred to as a so-called *shrinkage* method ((*cf.* for example Hastie *et al.*, 2009, pp. 61-79) and (Goodfellow *et al.*, 2016)).

We can further generalise linear regression to the classification scenario. First, note the different types of classification schemes. Consider \mathcal{G} , the discrete set of values which may be assumed by G , where G is used to denote a categorical output variable (instead of Y). Let $|\mathcal{G}| = K$ denote the number of discrete categories in the set \mathcal{G} . The simplest form of classification is known as binary classification and refers to scenarios where the input is associated with only one of two possible classes, *i.e.* $K = 2$. When $K > 2$, the task is known as multiclass classification. In multi-label classification an input may be associated with multiple classes (out of K available classes), where the number

of classes that each observation belongs to, is unknown. Here we start by introducing the two single label classification setups, *viz.* binary and multiclass classification.

In multiclass classification, given the input values \mathbf{X} , we would like to accurately predict the output, G , which we denote by \hat{G} . One approach would be to represent G by an indicator vector $\mathbf{Y}_G : K \times 1$, with elements all zero except in the G -th position, where it is assigned a 1, *i.e.* $Y_k = 1$ for $k = G$ and $Y_k = 0$ for $k \neq G$, $k = 1, 2, \dots, K$. We may then treat each of the elements in \mathbf{Y}_G as quantitative outputs, and predict values for them, denoted by $\hat{\mathbf{Y}} = [\hat{Y}_1, \dots, \hat{Y}_K]$. The class with the highest predicted value will then be the final categorical prediction of the classifier, *i.e.* $\hat{G} = \arg \max_{k \in \{1, \dots, K\}} \hat{Y}_k$.

Within the above framework we therefore seek a function of the inputs which is able to produce accurate predictions of the class scores, *i.e.*

$$\hat{Y}_k = \hat{f}_k(\mathbf{X}),$$

for $k = 1, \dots, K$. Here \hat{f}_k is an estimate of the true function, f_k , which is meant to capture the relationship between the inputs and output of class k . As with the linear regression case described above, we can use a linear model $\hat{f}_k(\mathbf{X}) = \hat{\mathbf{w}}_k^T \mathbf{X}$ to approximate the true function. The linear model for classification divides the input space into a collection of regions labelled according to the classification, where the division is done by linear *decision boundaries* (see Figure 1.2 for an illustration). The decision boundary between classes k and l is the set of points for which $\hat{f}_k(\mathbf{x}) = \hat{f}_l(\mathbf{x})$. These set of points form an affine set or hyperplane in the input space.

After the weights are estimated from the data, an observation represented by \mathbf{x} (including the unit element) can be classified as follows:

- Compute $\hat{f}_k(\mathbf{x}) = \hat{\mathbf{w}}_k^T \mathbf{x}$ for all $k = 1, \dots, K$.
- Identify the largest component and classify to the corresponding class, *i.e.* $\hat{G} = \arg \max_{k \in \{1, \dots, K\}} \hat{f}_k(\mathbf{x})$.

One may view the predicted class scores as estimates of the conditional class probabilities (or posterior probabilities), *i.e.* $P(G = k | \mathbf{X} = \mathbf{x}) \approx \hat{f}_k(\mathbf{x})$. However, these values are not the best estimates of posterior probabilities. Although the values sum to 1, they do not lie in the interval [0,1]. A way to overcome this problem is to estimate posterior probabilities using the *logit*




Figure 1.2: Linear model on simple binary classification dataset.

transform of $\hat{f}_k(\mathbf{x})$. That is,

$$P(G = k | \mathbf{X} = \mathbf{x}) \approx \frac{e^{\hat{f}_k(\mathbf{x})}}{\sum_{l=1} e^{\hat{f}_l(\mathbf{x})}}.$$

Through this transformation, the estimates of the posterior probabilities both sum to 1 and are contained in [0,1]. The above model is the well-known *logistic regression* model (Hastie *et al.*, 2009, p. 119). With this formulation there is no closed form solution for the weights. Instead, the weight estimates may be searched for by maximising the log-likelihood function. One way of doing this is by minimising the negative log-likelihood using gradient descent, which will be discussed in the next chapter.

Finally in this section, note that any supervised learning problem can also be viewed as a function approximation problem. Suppose we are trying to predict a variable Y given an input vector \mathbf{X} , where we assume the true relationship between them to be given by

$$Y = f(\mathbf{X}) + \epsilon,$$

where ϵ represents the part of Y that is not predictable from \mathbf{X} , because of, for example, incomplete features or noise present in the labels. Then in function approximation we are estimating f with an estimate \hat{f} . In parametric

function approximation, for example in linear regression, estimation of $f(\mathbf{X}, \theta)$ is equivalent to estimating the optimal set of weights, $\hat{\theta}$. In the remainder of the thesis, we refer to \hat{f} as the *model*, *classifier* or *learner*.

1.5 Outline

This chapter provided the context, motivation, objectives and theoretical background of this study. An outline of the remainder of the thesis follows below:

In Chapter 2, the theory underlying neural networks is described. The building blocks of neural networks are discussed, thereby introducing neurons, basic layers and the way in which neural networks are trained. The important concept of regularisation is also discussed. Using the perspective of representation- and manifold learning, we then attempt to gain insight into what happens inside an neural network.

Chapter 3 continues the discussion by focusing on the key advances in neural networks in recent times. The idea is that all concepts introduced in this chapter should potentially be able to facilitate the construction of improved deep neural networks on tabular data. Improved ways of fighting overfitting, such as data augmentation, the use of dropout and transfer learning, as well as the SotA training policy called *1Cycle* are analysed here. New developments in architectural design are also highlighted. The chapter concludes with approaches towards interpreting neural networks and their predictions.

Chapter 4 may be viewed as a core chapter of this thesis. It mainly serves as a literature review of all research with regard to deep learning for tabular data. The chapter is organised according to the modelling challenges faced when using deep learning for tabular data, investigating and comparing what other researchers have done in order to overcome these challenges. It will be seen that the key concept involves finding the right representation for tabular data. This may be done through embeddings, and by means of designing architectures that can efficiently learn feature interactions. This is for example done with attention models, possibly with the help of unsupervised pretraining.

In Chapter 5 we empirically evaluate several claims made in the literature. The aim of the chapter is to evaluate and compare different approaches to tackling the various challenges. Hence the main experiments involve evaluating neural networks at various samples sizes, evaluating potential gains from

doing unsupervised pretraining and using data augmentation, and comparing attention modules with classic fully-connected layers. We also illustrate a way in which resulting neural networks may be interpreted by means of permutation importance and knowledge distillation.

The thesis concludes in Chapter 6, where we summarise the work, some highlights and the main take-home points. The limitations of this study are discussed, and promising future research directions are identified.

Chapter 2

Neural Networks

2.1 Introduction

Not unlike most supervised machine learning models, a neural network (NN) is a function which maps inputs to outputs, *i.e.* $f : \mathbf{x} \rightarrow y$. The structure of f is often loosely compared to the structure of the human brain. Oversimplified, the brain consists of a collection of interconnected neurons. Each neuron can generate and receive signals. A received signal may be described as an input to a neuron, whereas a sent signal may be described as an output from that neuron. If two neurons are connected, it means that the output from the one neuron serves as input to the other. In a very simple model of the brain, one may argue that a neuron receives several signals, which it weighs and combines, and if the combined value of the inputs is higher than a certain threshold, the neuron sends a output signal to the next neuron. Figure 2.1 (a) provides a schematic of a biological neuron.

An artifical neural network tries to mimic this model of the human brain - it is set up to consist of several layers of connected units (or neurons). With exception of units in the first and final layers, each unit outputs a weighted combination of its inputs, combined with a simple non-linear transformation. In each layer of the neural network, the input is passed through each of the neurons. In turn, their output is passed to the next layer.

The transformation at each neuron is controlled by a set of parameters, also known as weights. Training a neural network involves tuning these weights in order to obtain some desired output. During training, the neural network receives as input a set of training data. The neural network weights are then

learned in such a way that, when given a new set of inputs, the output predicted by the neural network matches the corresponding response of interest as closely as possible. The process of using the training data to tweak the weights is done by means of an optimisation algorithm called Stochastic Gradient Descent (SGD).

Although recently there has been plenty of excitement around neural networks, it is well known that they were invented many years ago. The development of neural networks dates back at least as far as the invention of perceptrons in (Rosenblatt, 1962). It is also interesting to compare modern neural networks with the Projection Pursuit Regression algorithm in statistics (Friedman and Stuetzle, 1981). Only recently a series of breakthroughs allowed neural networks to be more effective, leading to the renewed interest in the field.

The aim of this chapter is to provide an overview of neural networks, emphasising the basic structure (§2.2) and the way in which they are trained (§2.3). This is done with a view to discuss modern neural network structures and training policies in Chapter 3, which in turn will help us shed light on Deep Learning for tabular data. Finally, the chapter concludes with a section on representation and manifold learning (§2.4) in an attempt to understand what a neural network is actually doing.

2.2 The Structure of a Neural Network

2.2.1 Neurons and Layers

In basic terms, a neural network processes an input \boldsymbol{x} by sending it through a series of layers. The neurons in each layer apply some transformation to their inputs, resulting in a set of outputs which are again passed on to the next layer of neurons. Eventually, the final layer produces the neural network output. In this section we provide more detail regarding the neural network structure. We start with a description of the operations inside each neuron, and follow with a discussion of the way in which the neurons may be connected in layers in order to form a complete neural network structure. Our discussion is based upon a simple regression example.

Suppose we are in pursuit of a function which is able to estimate some continuous target, y , given a p -dimensional input \boldsymbol{x} , *e.g.* estimating the taxi fare

from features such as distance travelled, time elapsed and number of passengers. A single neuron may act as such a function. It models y by computing a weighted average of the input features. This operation is illustrated in Figure 2.1 (b).



Figure 2.1: Comparison of a biological (a) and an artificial (b) neuron¹.

In equation form, this function can be written as:

$$w_1 \cdot x_1 + w_2 \cdot x_2 + \cdots + w_p \cdot x_p + b = y,$$

where $\{w_k\}_{k=1}^p$, are the weights applied to each of the inputs $\{x_k\}_{k=1}^p$ and b the constant bias term. Clearly, this equation is simply the very common linear model and thus also can be written as:

$$\mathbf{w}^\top \mathbf{x} + b = y,$$

where $\mathbf{x} = [x_1 \ x_2 \ \dots \ x_p]^\top$ is the input, $\mathbf{w} = [w_1 \ w_2 \ \dots \ w_p]^\top$ the weights and y the output. We may compress the above equation to $\mathbf{w}^\top \mathbf{x} = y$, where \mathbf{x} includes the bias term and the weight vector \mathbf{w} a unit element, i.e. $\mathbf{x} = [1 \ x_1 \ \dots \ x_p]^\top$ is the input, $\mathbf{w} = [b \ w_1 \ \dots \ w_p]^\top$.

The weights convey the importance of each input features in predicting the target. The larger $|w_k|$ is, the greater is the contribution of x_k towards the output. If $w_k = 0$, x_k has no influence on the target. However the weights are unknown and therefore we need to estimate them.

In linear regression this is done by means of the method of ordinary least squares. Since a neural network consists of many inter-connected neurons, an alternative estimation procedure is required. This is the topic of the next section.

Often a linear model will be too rigid to model a certain response of interest. In order to fit a more flexible model, we may add more neurons. Consider the

¹Image credit: <https://www.jeremyjordan.me/intro-to-neural-networks/>

use of two neurons, where the second neuron accepts the same input as the first neuron, but uses a different set of weights. Thus we have two different outputs produced by the two neurons, *i.e.* $z_1 = \mathbf{w}_1^\top \mathbf{x}$ and $z_2 = \mathbf{w}_2^\top \mathbf{x}$. In order to produce a final estimate from the initial two estimates, *viz.* z_1 and z_2 , they are passed to a third neuron. That is, $y = \mathbf{w}_3^\top \mathbf{z}$, where $\mathbf{z} = [z_1 \ z_2]^\top$. Figure 2.2 illustrates this pipeline in network form.




Figure 2.2: A simple neural network accepting p -sized inputs, with one hidden layer which has two neurons.

The first two neurons each received all p inputs and each produced a single output. These two outputs were received by the third neuron, and combined in order to produce the final output, *viz.* y . The operations performed by the first two neurons may be expressed as $\mathbf{z} = W\mathbf{x}^\top$, where

$$W = \begin{bmatrix} \mathbf{w}_1^\top \\ \mathbf{w}_2^\top \end{bmatrix} = \begin{bmatrix} w_{10} & w_{11} & w_{12} & \dots & w_{1p} \\ w_{20} & w_{21} & w_{22} & \dots & w_{2p} \end{bmatrix} \quad \text{and} \quad \mathbf{z} = [z_1 \ z_2]^\top.$$

The collection of these two neurons is what is called a layer. Since our third neuron (which is also a layer but with a single neuron) receives the output of this layer as input, it is possible to express the complete input-output relationship in one equation, *i.e.*

$$y = \mathbf{w}_3^\top \mathbf{z} = \mathbf{w}_3^\top W\mathbf{x}.$$

Note here that the weights from the first layer, W , and the third neuron, \mathbf{w}_3 , can be collapsed into a single vector \mathbf{w} , effectively reducing all of the neuron operations back into a single neuron representation and thus the fitted model is still linear. In order to fit a non-linear model, a non-linear transformation

function is applied to the output of each layer. This function is called an *activation function*.

Subsequently the neural network equation can be written as

$$y = a_2(\mathbf{w}_3^T a_1(W\mathbf{x})) ,$$

where a_1 denotes the activation function applied after the first (linear) layer, and where a_2 is the activation function applied after the final layer.

The introduction of non-linear activation functions serves to enlarge the class of functions that can be approximated by the network *i.e.* making it possible for the network to learn complex, non-linear relationships between the inputs and the outputs. Next, we briefly discuss the various activation functions.

2.2.2 Activation Functions

There are plenty of activation functions to choose from, since any simple non-linear and differentiable function can be used. Originally, the *sigmoid* activation function was a common choice (Rumelhart *et al.*, 1988). It can be expressed as $\text{sigmoid}(x) = \frac{1}{1+e^{-x}}$. The shape of the sigmoid activation can be seen in Figure 2.3 (a). It has an S-shape and its values range between 0 and 1. The reason it fell out of favour is because of issues related to the gradient based optimisation procedure of neural networks (which we discuss in more detail in Section 2.3). One of the issues is that the values of sigmoid activations are not centered around zero. This may cause the gradient weight updates to veer too far in different directions.

The hyperbolic tangent or the *tanh* activation function, on the other hand, returns outputs centered around zero. It is expressed as $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ and its shape can be seen in Figure 2.3 (a). But the problem with both the sigmoid and tanh activation functions is that they can cause the gradients to saturate during training. By observing the tails of the sigmoid and tanh functions, it is clear that the gradient tends to zero as $|x| \rightarrow \infty$. This can cause the weight updates to be very close to zero and thus resulting in the network getting stuck at a certain point in the parameter space. Furthermore, the maximum gradient of the sigmoid activation function is 0.25 (at $x = 0.5$), and because of the nature of the chain rule this causes the lower layers to train much slower than the higher layers. The tanh activation function typically have larger derivatives




Figure 2.3: Plots of various activation functions (a) and their local derivatives (b).

than the sigmoid and thus it is not as susceptible to this vanishing gradient problem, however, it is still not immune to it. The local derivatives of the activation functions discussed here are plotted in Figure 2.3 (b). This will become more obvious after Section 2.3.

To date, the most popular choice in activation function is the *Rectified Linear Units* (ReLU) non-linearity. It is defined as:

$$\text{relu}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}.$$

Again, its shape and derivative is plotted in Figure 2.3 (a) and (b), respectively. The ReLU limits the gradient vanishing problem as its derivative is always 1 when x is positive. This results in significantly shorter steps to convergence, as found by the authors in (Krizhevsky *et al.*, 2012). However, ReLUs may suffer from the “dead ReLU” problem, which is, if a ReLU neuron gets clamped to zero, then its weights will get zero gradient and may remain permanently “dead” during training. This sparsity of the activations is what some believe is the reason for the effectiveness of ReLUs (Sun *et al.*, 2014). If dead activations want to be avoided, alternative activation functions are PReLU (He *et al.*, 2015a) and Leaky ReLUs (Maas *et al.*), but they fall out of the scope of this work.

Typically, choosing the right activation function for a specific task comes down to trial and error, however, in most cases it would be sufficient to use ReLUs after each hidden layer. When doing classification, it is often useful to produce outputs between 0 and 1 for each category as estimates of the

conditional class probabilities. Therefore we may use a sigmoid activation function on the output layer. When doing (single label) multiclass classification it is also desirable for these outputs to sum to 1, similar to probabilities. In that case we may use an activation function called the *softmax*, which is exactly the logit transformation introduced in §1.4: $\text{softmax}(\mathbf{x})_k = \frac{e^{x_k}}{\sum_{l=1}^K x_l}$. For regression tasks, we mostly omit an activation function on the output layer.

In our experiments we will show how the different activation functions perform on tabular data and compare it to the *Scaled Exponential Linear Units* (SELU) (Klambauer *et al.*, 2017) activation, which is supposed to allow us to more effectively train deeper neural networks. The neural network depth refers to its number of hidden layers, whereas the width of a layer refers to the number of neurons it consists of. The choice in network depth and layer width is the topic of the next section.

2.2.3 Size of the Network

The network depth and the width of its hidden layers (*i.e.* the size of the network) are hyperparameters of the model. They control the capacity or the flexibility of a neural network; referring to its ability to model complex functions. The bigger the network, the more flexible the model is. Like with most such parameters, increasing the model size is usually beneficial up to a certain point, after which its test performance starts to degrade (overfitting). In addition, more layers and more neurons adds to the time taken and hardware required to train the neural network.

Thus the challenge is to find a network size that is large enough to capture all the complexities in the data, but small enough to avoid overfitting. Currently the best way of finding the optimal size of a network for a given problem is by experimentation. Note, that this is true for many of the components of neural networks and deep learning. Later in this chapter, §2.3.5, we will see why tuning the network size is not necessarily the best way of controlling overfitting.

Theoretically, according to the universal approximation theorem (Cybenko, 1989), a neural network with a single hidden layer and a finite number of neurons can approximate any continuous function. Which begs the question as to why we need more hidden layers. Even though a neural network can represent any function, it does not mean that the available learning algorithms

can find these optimal weights (Ba and Caurana, 2013). It may also be that the number of neurons needed for a single hidden layer network to represent a specific function is infeasibly large. By choosing deeper networks we are assuming that the function we are trying to learn is composed of several simpler functions. Building in this prior proves to be useful empirically (Goodfellow *et al.*, 2016, pp.197-198), especially for tasks of a hierarchical nature like in computer vision.

In our experiments we will see what effect the layer width and network depth have on the generalisation performance on different tabular datasets. The networks used for tabular data learning problems are typically much shallower than those used for say computer vision and NLP. We investigate why this is the case. Later in this chapter we view this problem from a representation learning perspective which will also give us some insight on the matter. But first, in the next section, we discuss the process of training a neural network.

2.3 Training a Neural Network

2.3.1 Optimisation

Statistical learning and optimisation are closely related. Optimisation refers to the task of either minimising or maximising some function $J(x)$ by altering x . The function we want to optimise is called the *objective function*. When we are minimising the objective function, we may also refer to the objective function as the *cost* or *loss function*. These terms will be used interchangeably throughout the thesis.

As mentioned in the previous chapter, parameter estimation (or optimisation) of a linear (or logistic regression) model is usually done using OLS or maximum likelihood estimation (MLE). In this section, however, we discuss an alternative parameter estimation method which is also relevant for the optimisation of neural networks.

Consider the MSE loss function:

$$\begin{aligned} L &= \sum_{i=1}^N L_i \\ &= \sum_{i=1}^N \sum_{k=1}^K (y_{ik} - f_k(\mathbf{x}_i))^2 \\ &= \sum_{i=1}^N \sum_{k=1}^K (y_{ik} - \mathbf{w}_k^T \mathbf{x}_i)^2, \end{aligned}$$

where $f_k(\cdot)$ in this case is the linear model used to predict the k -th class posterior probability. Although the MSE loss is mostly used in a regression setup and not really well suited for classification, we make use of it here for illustration purposes.

To find the weights, \mathbf{w} , that minimise L , we can follow a process of iterative refinement. That is, starting with a random initialisation of \mathbf{w} , one iteratively updates the values such that L decreases. The updating steps are repeated until the loss converges. In order to minimise L with respect to \mathbf{w} , we calculate the gradient of the loss function at the point $L(\mathbf{x}; \mathbf{w})$. The gradient (or slope) of the loss function indicates the direction in which the function has the steepest rate of increase. Therefore, once we have determined this direction, we can update the weights by a step in the opposite direction - thereby reaching a smaller value of L .

The gradient of L_i is computed by obtaining the partial derivative of L_i with respect to \mathbf{w}_k , *i.e.*:

$$\frac{\partial L_i}{\partial \mathbf{w}_k} = -2(y_{ik} - \mathbf{w}_k^T \mathbf{x}_i) \mathbf{x}_i.$$

After obtaining the above N partial derivatives, an update at the $(r+1)$ -th iteration may be obtained as follows:

$$\mathbf{w}_k^{(r+1)} = \mathbf{w}_k^{(r)} - \gamma \sum_{i=1}^n \frac{\partial L_i}{\partial \mathbf{w}_k^{(r)}},$$

where γ is called the *learning rate* and determines the size of the step taken towards the optimal direction. One typically would like to set the learning rate small enough so that one does not overshoot the minimum, but large enough to limit the number of iterations before convergence. The learning rate is a crucial parameter when training neural networks; we will discuss its significance in §2.3.6.

The procedure of repeatedly evaluating the gradient of the objective function and then performing a parameter update, is called *gradient descent* (Cauchy, 1847). Gradient descent forms the basis of the optimisation procedure for neural networks.

Note that a weight update is made by evaluating the gradient over a set of observations, $\{\mathbf{x}_i, i = 1, \dots, n\}$. One of the advantages of gradient descent is that at an iteration, the gradient need not be computed over the complete training dataset, *i.e.* $n \leq N$. When updates are iteratively determined by using subsets of the data, the process is called *mini-batch gradient descent*. This is extremely helpful in large-scale applications, since it obviates computation of the full loss function over the entire dataset. This leads to faster convergence, because of more frequent parameter updates, and allows processing of data sets that are too large to fit into a computer's memory. The choice regarding batch size depends on the available computation power. Typically a batch consists of 64, 128 or 256 data points, since in practice many vectorised operation implementations work faster when their inputs are sized in powers of 2. The gradient obtained using mini-batches is only an approximation of the gradient of the full loss but it seems to be sufficient in practice (Li *et al.*, 2014). Note at this point that the collection of iterations needed to make one sweep through the training data set is called an *epoch*.

The extreme case of mini-batch gradient descent is when the batch size is selected to be 1. This is called *Stochastic Gradient Descent* (SGD). Recently SGD has been used much less, since it is more efficient to calculate the gradient in larger batches compared to only using one example. However, note that it remains common to use the term SGD when actually referring to mini-batch gradient descent. Gradient descent in general has often been regarded as slow or unreliable but it works well for optimising neural networks. SGD will most probably not find even a local minimum of the objective function. It typically however finds a very low value of the cost function quickly enough to be useful.

2.3.2 Optimisation Example

To illustrate the SGD algorithm, consider the linear model in a classification context. Suppose we are given a training dataset with two-dimensional inputs and only two possible classes. Let the data be generated in the same way as described in (Hastie *et al.*, 2009, pp. 16-17).

We want to fit a linear regression model to the data such that we can classify an observation to the class with the highest predicted score. In the binary case it is only necessary to model one class probability and then assign an observation to that class if the score exceeds some threshold (usually 0.5), otherwise it is assigned to the other class. Therefore the decision boundary is given by $\{\mathbf{x} : \mathbf{x}^T \hat{\mathbf{w}} = 0.5\}$.

The example is illustrated in Figure 2.4. The colour shaded regions represent the parts of the input space classified to the respective classes, as determined by the decision boundary based upon OLS parameter estimates. Gradient descent was applied to the determine the optimal weights using a learning rate of 0.001. Since the total number of training observations are small, it is not necessary to use SGD. In Figure 2.4, the dashed lines represent the decision boundary defined by the gradient descent parameter estimates at different iterations. We observe that initially the estimated decision boundary is far from the OLS solution, but as the update iterations proceed, the decision boundary is rotated and translated until finally matching the OLS line. It took 29 iterations for the procedure to reach convergence.

2.3.3 Backpropogation

In Section 2.3.1 we discussed how to fit a linear model using the Stochastic Gradient Descent optimisation procedure. Currenlty, SGD is the most effective way of training deep networks. To recap, SGD optimises the parameters θ of a network to minimise the loss,

$$\theta = \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N l(\mathbf{x}_i, \theta).$$

With SGD the training proceeds in steps and at each step we consider a mini-batch of size $n \leq N$ training samples. The mini-batch is used to approximate the gradient of the loss function with respect to the paramaters by computing,

$$\frac{1}{n} \frac{\partial l(\mathbf{x}_i, \theta)}{\partial \theta}.$$

Using a mini-batch of samples instead of one at a time produces a better estimate of the gradient over the full training set and it is computationally much more efficient.




Figure 2.4: Plots of the gradient descent example. (a) The data points in input space. The shades in the background represent the class division in input space, with the decision boundary determined by linear least squares estimation. The dashed lines represent the decision boundaries learned at different iterations. (b) The loss calculated at each iteration.

This section discusses the same procedure, but applied to a simple single hidden layer neural network for multiclass classification, decomposed as:

$$\begin{aligned} f_k(\mathbf{x}) &= g_k(\boldsymbol{\beta}_k^\top \mathbf{z}), \quad k = 1, \dots, K \\ z_m &= \sigma(\boldsymbol{\alpha}_m^\top \mathbf{x}), \quad m = 1, \dots, M \end{aligned}$$

where $\sigma(\cdot)$ is the sigmoid activation and $g(\cdot)$ the softmax activation. The neural network has a set of unknown adjustable weights that defines the input-output function of the network. They are the parameters of the linear function of the inputs, $\boldsymbol{\alpha}_m = (\alpha_{0m}, \alpha_{1m}, \dots, \alpha_{pm})$, and the parameters of the linear transformation of the derived features, $\boldsymbol{\beta}_k = (\beta_{0k}, \beta_{1k}, \dots, \beta_{mk})$. Denote the complete set of parameters by θ . Then the objective function for regression can be chosen as the sum-of-squared-errors:

$$L(\theta) = \sum_{k=1}^K \sum_{i=1}^N (y_{ik} - f_k(\mathbf{x}_i))^2$$

and for classification, the cross-entropy:

$$L(\theta) = - \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log f_k(\mathbf{x}_i),$$

with corresponding classifier $G(\mathbf{x}) = \arg \max_k f_k(\mathbf{x})$. Since the neural network for classification is a linear logistic regression model in the hidden units, the parameters can be estimated by maximum likelihood. According to Hastie *et al.* (2009, p. 395), the global minimiser of $L(\theta)$ is most likely an overfit solution and we instead require regularisation techniques when minimising $L(\theta)$. Furthermore, as the network becomes larger, MLE becomes intractable.

Therefore, one rather uses gradient descent and the *backpropagation* algorithm (Rumelhart *et al.*, 1988) to minimise $L(\theta)$. This is possible because of the modular nature of a neural network, allowing the gradients to be derived by iterative application of the chain rule for differentiation. This is done by a forward and backward sweep over the network, keeping track only of quantities local to each unit.

In detail, the backpropagation algorithm for the sum-of-squared error objective function,

$$\begin{aligned} L(\theta) &= \sum_{i=1}^N L_i \\ &= \sum_{i=1}^N \sum_{k=1}^K (y_{ik} - f_k(\mathbf{x}_i))^2, \end{aligned}$$

is as follows. Following the chain-rul, the relevant derivatives for gradient descent are:

$$\begin{aligned} \frac{\partial L_i}{\partial \beta_{km}} &= -2(y_{ik} - f_k(\mathbf{x}_i))g'_k(\boldsymbol{\beta}_k^T \mathbf{z}_i)z_{mi}, \\ \frac{\partial L_i}{\partial \alpha_{ml}} &= - \sum_{k=1}^K 2(y_{ik} - f_k(\mathbf{x}_i))g'_k(\boldsymbol{\beta}_k^T \mathbf{z}_i)\beta_{km}\sigma'(\boldsymbol{\alpha}_m^T \mathbf{x}_i)x_{il}. \end{aligned}$$

Given these derivatives, a gradient descent update at the $(r+1)$ -th iteration has the form,

$$\begin{aligned} \beta_{km}^{(r+1)} &= \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial L_i}{\partial \beta_{km}^{(r)}}, \\ \alpha_{ml}^{(r+1)} &= \alpha_{ml}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial L_i}{\partial \alpha_{ml}^{(r)}}. \end{aligned}$$

Now write the gradients as

$$\begin{aligned}\frac{\partial L_i}{\partial \beta_{km}} &= \delta_{ki} z_{mi}, \\ \frac{\partial L_i}{\partial \alpha_{ml}} &= s_{mi} x_{il}.\end{aligned}$$

The quantities, δ_{ki} and s_{mi} are errors from the current model at the output and hidden layer units respectively. From their definitions, they satify the following,

$$s_{mi} = \sigma'(\boldsymbol{\alpha}_m^T \mathbf{x}_i) \sum_{k=1}^K \beta_{km} \delta_{ki},$$

which is known as the backpropogation equations. Using this, the weight updates can be made with an algortihm consisting of a forward and a backward pass over the network. In the forward pass, the current weights are fixed and the predicted values $\hat{f}_k(\mathbf{x}_i)$ are computed. In the backward pass, the errors δ_{ki} are computed, and then backpropogated via the backpropogation equations to obtain s_{mi} . These are then used to update the weights. This approach extends naturally to any size network and differentiable layers (functions).

Backpropogation is simple and its local nature (each hidden unit passes only information to and from its connected units) allows it to be impleted efficiently in parallel. The other advantage is that the computation of the gradient can be done on a batch (subset of the training set) of observations. This allows the network to be trained on very large datasets.

In summary, training a neural network consists of these four steps:

1. Initialise the network weights: give random set of numbers to the network parameters.
2. Forward propogation: pass the input through the network layers to produce an output.
3. Calculate the error: compare the predicted output with the true output and measure the difference using an objective function.
4. Backward propogation: calculate the gradients of the objective function with respect to the weights and update the weights accordingly.

These four steps are typically repeated until convergence of the loss function. It can take many training epochs for the objective function to converge.

2.3.4 Weight initialisation

Before one begins training a neural network, one needs to initialise its weights. We expect a well trained neural network to have an equal number of weights smaller than zero and greater than zero. If we initialise all weights to be zero, every neuron will compute the same output and therefore produce the same gradient and undergo the same weight update. We still want to initialise the weights as small as possible but each one unique. Sampling from zero mean and unit variance Gaussian distribution is a natural fit, however, the variance of the outputs from a randomly initialised neuron grows with increasing number of inputs. Therefore we may scale the weight vector by the square root of its number of inputs to normalise the variance to 1. This ensures that all neurons in the network initially have approximately the same output distribution and empirically improves the rate of convergence. (He *et al.*, 2015b) derived a weight initialisation specifically for ReLU neurons.

2.3.5 Basic Regularisation

From §2.2.3 it seems that smaller neural networks can be preferred if the data is not complex enough to prevent overfitting. However, this is untrue and there are much better ways of regularising neural networks, which we discuss shortly and in the next chapter. Smaller networks are harder to train with local methods such as gradient descent because their loss functions have poor local minima that are easy to converge to. Whereas the local minima of larger networks are much better in terms of their actual loss (Choromanska *et al.*, 2014).

One of the preferred ways of fighting overfitting in neural networks is by using L1 or L2 regularisation, *i.e.* adding a penalty term proportional to the magnitude of the weights to the objective function. This is exactly what is done in Ridge Regression and the Lasso (Hastie *et al.*, 2009, Ch. 4). In L2 regularisation we add the term $\frac{1}{2}\lambda w^2$ to the objective and in L1 regularisation the penalty term $\lambda|w|$. These penalties force the weights to be small and the λ parameter controls the strength of the regularisation. The “ $\frac{1}{2}$ ” in front of the L2 penalty is added for convenience to make the derivative equal to λw . This makes it then equivalent to decaying each of the weights linearly towards zero: $w' = w - \lambda w$, which is also known as *weight decay*. In practice, L2 regularisation outperforms L1 regularisation.

Another basic way of preventing the neural network to overfit is to stop the training process “early” and to not train until the training loss converges. A converged training loss is not equivalent to an optimal test loss and therefore the test or validation loss should be observed during training. One approach is to stop training when the validation loss stops decreasing. As mentioned before, the learning rate also plays a big part in finding the optimal weights. Next we discuss how we can tune the learning rate to train faster and to find better local minima. More advanced regularisation techniques are discussed in Chapter 3.

2.3.6 Adaptive Learning Rates and Annealing

Although the right learning rate can reduce training time and improve performance, there is no silver bullet when it comes to setting the parameter. A small learning rate slows down the training time, but is safer against overfitting and overshooting the optimal solution. With a large learning rate, convergence may be reached quicker, but the optimal solution may not have been found. One could do a line search of a range of possible values, but this usually takes too long for bigger networks.

However, one does not need to keep the learning rate fixed throughout the training process. A popular approach is to decrease the learning rate by a fraction after a certain number of epochs or as the validation loss starts to converge, as done in (He *et al.*, 2015a) for example. The intuition is that larger steps can be taken when we are still far away from an optimal position on the loss surface, and then gradually take smaller steps once we get closer to not overshoot it. Note, that this means that we would also need to tune the rate of decrease and the time steps of each decrease during training. Fortunately, it is believed that the learning algorithms are not that sensitive to this choice.

In addition, there are ways of manipulating the learning rates at a local level, as opposed the aforementioned global methods. Adagrad (Duchi *et al.*, 2011) is an adaptive learning rate method which increases the learning rate at neurons with small gradients and *vice versa*. Adam (Kingma and Ba, 2014) is the most commonly used weight update approach. It also uses the magnitude of the gradient to control each weight update, in addition to the previous iteration’s gradients and it combines them in a smooth fashion. This resembles the physical property of momentum (Bengio *et al.*, 2012). For more detail, the

reader can refer to the cited publications as it falls out of the scope of this work.

2.4 Representation Learning

We are now familiar with the mathematical operations of basic layers, how they are connected and how their weights are tweaked to minimise a loss function. In this section we will discuss why this works and what the neural network is actually doing to model the data. The central idea is that of a data *representation* (Bengio *et al.*, 2013) and that at each layer of the network the data is transformed into a higher-level abstraction of itself. Understanding and interpreting neural networks remains a challenge (Frosst and Hinton, 2017), but the notion of learning an optimal data representation allows us to gain a deeper intuition of the inner mechanics of neural networks.

Machine learning models are very sensitive to the form and the properties of the input given to it. Thus a large part of building machine learning models is to find the best way of representing the raw data to make it easier for the models to extract useful information. This is typically a laborious manual task of creating, analyzing, evaluating and selecting appropriate features² requires practitioner expertise and domain knowledge. This *feature engineering* process is more trial-and-error than a systematic recipe. Therefore if one can effectively automate this process, it will save a lot of time and raise the performance ceiling of models. Automatically learning representations of the data that make it easier to extract useful information for classifiers or other predictors is called representation learning (Bengio *et al.*, 2013).

A neural network can be viewed from the perspective of representation learning. Consider a classification task. Since the final layer of a neural network is a linear model, in order for the network to produce accurate predictions, the previous layers should be able to project the data into a space where the classes are linearly separable. Thus the network learns a representation of the data that is optimal for classification.

Each of the simple but non-linear modules of a neural network transforms the representation at one level (starting with the raw input) into a representation at a higher, slightly more abstract level. With the composition of enough such transformations, very complex functions can be learned (Lecun *et al.*, 2015).

²<http://blog.kaggle.com/2014/08/01/learning-from-the-best/>

These transformations can emphasise (and create) features that are important for discrimination and drop those which are redundant.

Let us work through a simple example to illustrate the representations learned by a neural network. Consider a dataset with two classes; the two curves on a plane shown in Figure 2.5. Clearly, the observations from the two classes are not linearly separable in their raw form. Thus if we fit a single layer neural network (*i.e.* only an output layer) to this data, we will get an unsatisfactory decision boundary, since the decision boundary can only be linear, as shown in Figure 2.6. However, if we fit a two-layer neural network, where the hidden layer has two neurons and a sigmoid activation to the same dataset, the decision boundary perfectly separates the two classes. This is shown in Figure 2.7.




Figure 2.5: Simple dataset with two linearly inseparable classes.

Since the hidden layer consists of only two neurons, we are able to plot the output from the hidden layer after the raw data has passed through it. This is depicted in Figure 2.8. This shows how the hidden layer projected the input data into a space where the observations from the two classes are linearly separable. Which leaves it to the final layer to find the best hyperplane between them.




Figure 2.6: Decision boundary of 1-layer neural network.




Figure 2.7: Decision boundary of 2-layer neural network.




Figure 2.8: Hidden representation of 2-layer neural network.

Of course this is a very simple example, but the same concepts apply to more complicated datasets and models. However, even though it is technically possible to separate any arrangement of points with a sufficiently large network³, in reality it can become quite challenging to find such representations. This is where the need for more data, regularisation, smarter optimisation procedures and architecture design arises. Without the aforementioned, it is likely that the network will get stuck in a sub-optimal local minima, not being able to find the optimal representation of the data. In the following chapters we explore the approaches available to find optimal representations of tabular data for classification and regression tasks.

³<http://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>

Chapter 3

Deep Learning

3.1 Introduction

Deep learning is a broadly used term. The difference between a classical neural network and a deep neural network, is merely its number of layers. Even a network with two hidden layers is sometimes referred to as deep.

In this work, we use the term deep learning to refer to modern developments in the space of neural networks. In the last decade, tons of contributions have been made to the field: new types of layers, different ways of training and novel approaches to fight overfitting. Reviewing these developments requires much more space than a chapter, therefore, for the purpose of this work, we only highlight the developments we feel are important for our next chapter on deep learning for tabular data.

Practically all deep learning developments are aimed at combatting overfitting, either explicitly or implicitly. Since neural networks can already approximate any function, but our learning algorithms cannot necessarily find these solutions, we need ways of learning more efficiently from the data. This chapter starts by introducing a different class of neural networks called *autoencoders* (§3.2). Autoencoders are mostly used as an unsupervised learning technique to help us learn more robust representations from data without labelled data, which we may then transfer to supervised learning algorithms. This process of transferring knowledge from one network to another is called transfer learning, which we discuss in §4.4.2.

Thereafter, we look at layers that help with regularisation: *dropout* (§3.4.1) and *batch normalisation* (§3.5.1). Data augmentation is a technique used to

artificially enlarge the training dataset and thus also helps fighting overfitting. We discuss this in §3.4.2. The other types of layers we look at in this chapter are, *skip-connections* (§3.5.2), *attention* modules (§3.5.4) and *embedding* layers (§3.5.3). They are designed to make deeper neural networks more effective, let them focus more on the important features and to process discrete input, respectively.

Then we devote a section to a training policy called the *1cycle* policy (§3.6). This method provides a way of automatically finding a good learning rate and drastically reduces the number of training iterations needed. We conclude this chapter with a discussion on the interpretability of neural networks (§4.5).

3.2 Denoising Autoencoders

The basic autoencoder is a neural network that is trained to attempt to reconstruct its inputs. The simplest form of an autoencoder is a neural network with one hidden layer and an output layer the same size of the input layer (illustrated in Figure 3.1). The linear layer transforming the input to the hidden layer is referred to as the encoder, $\mathbf{z} = f(\mathbf{x})$, and the layer producing the output from the hidden layer is called the decoder, $\mathbf{x}' = g(\mathbf{z})$. The autoencoder can be trained in the same way as any other neural network, but by minimising a reconstruction loss. A reconstruction loss measures the distance between the reconstruction of the input and the actual input and thus the MSE loss function is common choice for continuous data.

Autoencoders technically belong to the self- (or semi-) supervised class of methods, although many still think of it as unsupervised. It is unsupervised in the sense that it does not require labelling, but it is still supervised in the sense that it predicts an output.

In this setting, if the number of neurons in the hidden layer is greater than or equal to the number of input features, the autoencoder will be able to perfectly reconstruct \mathbf{x} from \mathbf{z} , *i.e.* $\mathbf{x}' = \mathbf{x}$. However this is not a very useful model and instead autoencoders are usually built with some type of constraint imposed. A common option is to restrict the number of neurons in the hidden layer to be smaller than the number of input features. This forces the autoencoder to capture only the most useful properties of the data in the hidden representation and can thus effectively be used as a way of dimensionality reduction (Hinton and Salakhutdinov, 2006). When there are no non-linear activation functions




Figure 3.1: A simple single hidden layer autoencoder with 4-dimensional inputs and 2 neurons in the hidden layer.

after the linear layers, one can show that this autoencoder is equivalent to a *Principal Component Analysis* (PCA) of the inputs.

There is no restriction in the number or the size of the layers used for the encoder and decoder, and if activation functions are used, one can potentially learn a more powerful non-linear generalisation of PCA. However, if the encoder and decoder networks are allowed too much capacity, the autoencoder can learn to reconstruct the inputs without learning useful properties of the data. It is hard to tell if an autoencoder has learned a useful latent representation of the data. One way to evaluate it is to use the features extracted by the encoder for a supervised learning task and compare its performance with a model using the raw data as inputs. Autoencoders have also been used to initialise the weights of a supervised learning network, *i.e.* using the learned weight of the autoencoder as the initial weights for the supervised learning network of the same size (Larochelle *et al.*, 2009).

An autoencoder can learn to perfectly copy the input if the encoder and decoder have enough capacity, even though the latent representation is of a smaller dimension than the inputs, and thus will not learn useful features of the data. Therefore we may want to consider other types of constraints to impose on the autoencoder.

A *denoising autoencoder* (DAE) (Vincent *et al.*, 2008) first adds noise to its inputs before it is passed to the encoder. Thus in order to minimise the reconstruction loss, the DAE should learn how to reconstruct the original input

from a corrupted version of itself. (Alain and Bengio, 2014) shows that DAEs do learn useful structures of the data.

The choice of the type of noise added to the inputs depends on the data types. One can block out inputs with zeros if zeros have no other meaning in the data or one might want to add gaussian noise to continuous outputs as long as it falls within the true range of the features. If there is too much noise in the corruption step, the DAE might not learn anything useful. The type and amount of noise to be use are things that should be experimented with.

Another type of autoencoder is a *Variational Autoencoder* (VAE) (Kingma and Welling, 2013). Instead of encoding the data into a latent vector representation, a VAE encodes the data as a Gaussian distribution. In the decoding step, observations are then sampled from the learned distribution before passing it to the fully connected layers. This is an interesting direction for future research but not in the scope of this work.

3.3 Pretraining and Transfer Learning

Unsupervised learning played a key part in the rise of deep learning (debatedly starting with (Hinton *et al.*, 2006)). It made it easier to train deeper neural networks with a process called *pretraining*. Pretraining is based on the assumption that representations learned from one task can sometimes be useful for accomplishing other tasks in the same input domain. For example, we can use a DAE to learn the structure of the data from unlabelled data and then use what it has learned as features or initialisation for a supervised learning task with the same type of inputs. We can also do supervised pretraining by first training the network to estimate one target variable and then use those weights or features when training to estimate a different target variable, for example.

Unsupervised pretraining for supervised learning is very common in NLP (Devlin *et al.*, 2018, Howard and Ruder (2018)) and supervised pretraining for supserivsed learning is very common in computer vision (Yosinski *et al.*, 2015, He *et al.* (2015a)). It is not yet clear theoretically why pretraining works. It may be that using the pretrained weights as initialisation to the supervised model provides a better starting position on the loss surface and thus have regulartory effects (Goodfellow *et al.*, 2016, Ch. 14). Pretraining is most effective when there is little data available for the supervised task but a lot of data available for the pretraining task.

The process of transferring what is learned doing one task (for example in pretraining) to another task (supervised learning for examples) is called *transfer learning*. When done effectively, pretraining and transfer learning can together dramatically reduce the number of training samples and computational power need to train a model.

(Zeiler and Fergus, 2014) provides an insightful investigation of the what a CNN learns when trained on image data. From observing the types of features extracted for a trained image model (Figure 3.2) it becomes clear why they are also effective on other image datasets. The learned filters seemed to look for generic image features like edges and color gradients which are useful for most computer vision tasks. (Howard and Ruder, 2018) describes a good practice when doing transfer learning via gradual unfreezing and discriminative fine-tuning.




Figure 3.2: Visualising the first layer convolutional filters learned by a neural network in a large image dataset.

3.4 More Regularisation

In §2.3.5 we discussed basic regularisation methods for neural networks. There are many other ways of combatting overfitting. The following section discusses two regularisation techniques proven to be extremely powerful in almost every application of deep learning.

3.4.1 Dropout

Dropout consists of randomly setting the output of a hidden neuron to zero with probability p . See Figure 3.3 for an illustration of how dropout effects the connections between neurons. The neurons which are set to zero do not contribute to the forward pass and do not participate in backpropogation. Every time an input is presented, the neural network samples a different set of neurons to be dropped out.

This technique ensures that a neuron does not rely on the signals of a particular set of other neurons, It is therefore forced to learn more robust features that are useful in conjunction with many different random subsets of the other units (Krizhevsky *et al.*, 2012).

In the original paper they used $p = 0.5$ but the optimal selection will depend on the use case and is typically found by experimentation. At test time, no units are dropped out and their output is multiplied by $1 - p$ to compensate for the fact that all of the units are now active. Dropout does tremendously well to combat overfitting, but it slows down the convergence time of training.

There are also parallels to be drawn between dropout and ensembling approaches (Hinton *et al.*, 2012). Since, at each training iteration a unique set of neurons are active which may then each be viewed as unique models. During training these models are combined - similar to the process of ensembling models.

3.4.2 Data Augmentation

Recall that our aim with predictive models is to generalise well to an unseen test set. In an ideal world we would train a model on all possible variations of the data to capture all interactions and relationships. This is not possible in the real world. Such a dataset is not available and would be infinitely large. And even if it were, machine learning will then be unnecessary since all the




Figure 3.3: The effect dropout has on connections between neurons.

possible observations are observable and one could simply use a lookup or a search when predicting outcomes.

In reality we have a finite subset of the full data distribution to train on. Any new samples with unique feature combinations will likely improve the models generalisability. If the collection of new samples is not available, we can try to artificially create more through *data augmentation*.

This is a standard approach especially in computer vision applications. For example, from a single image, we can rotate it, flip it horizontally, shift it in any direction, crop it, and many other transformations without destroying the semantic content of the image. But by doing so we are artificially increasing the size of the training set to help with overfitting. In Figure 3.4 we see how a single image of a cat is turned into eight images by doing random transformations of the original image. In all of these images it is still possible to tell that it is an image of a cat. Of course this is not as effective as genuine new data samples, but it is a very effective and efficient substitute (Perez and Wang, 2017). Data augmentation consistently leads to improved generalisation.



Figure 3.4: Example data augmentations for images.

Data augmentation can be formalised by the *Vicinal Risk Minimisation*

principle (Chapelle *et al.*, 2001) where human knowledge is required to describe a vicinity around each observation in the training data so that artificial examples can be drawn from the vicinity distribution of the training sample to enlarge it. In image classification one can define the vicinity of an image as the set of its horizontal reflections and minor rotations, for example.

Note that data augmentation is dataset and domain dependent and that the augmentations should preserve the semantic content or the signal in the observation. For example, taking too small crops of an image will ignore the context and may make it impossible to tell what is inside. And augmentations that are suitable for image data does not necessarily make sense for text data.

3.5 Modern Architectures

In the following sections we highlight some of the recently developed neural network layers and modules. These operations were designed to make training more robust and efficient, to help learn more robust representations and be able to model more useful features, among others.

3.5.1 Normalisation

One of the things that complicate the training of neural networks is the fact that hidden layers have to adapt to the continuously changing distribution of its inputs. The inputs to each layer are affected by the parameters of all its preceding layers and a small change in a preceding layer can lead to a much bigger difference in output as the network becomes deeper. When the input distribution to a learning system changes, it is said to experience covariate shift (Shimodaira, 2000).

Using ReLUs, careful weight initialisation and small learning rates can help a network to deal with the internal covariate shift. However, a more effective way would be to ensure that the distribution of non-linearity inputs remains more stable while training the network. (Ioffe and Szegedy, 2015) proposed *batch normalisation* to do just that.

A batch normalisation layer normalises its inputs to a fixed mean and variance (similar to how the inputs of the network are normalised) and therefore it can be applied before any hidden layer in a network to prevent internal covariate shift. The addition of this layer dramatically accelerates the training

of deep neural networks, also because it can be used with higher learning rates. It also helps with regularisation (Ioffe and Szegedy, 2015), therefore in some cases dropout is not necessary.

The batch normalising transform over a batch of univariate inputs, x_1, \dots, x_n is done by the following steps:

1. Calculate the mini-batch mean, μ , and variance, σ^2 :

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

2. Normalise the inputs,

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}},$$

where ϵ is a constant to ensure numerical stability.

3. Scale and shift the values,

$$y_i = \gamma \hat{x}_i + \beta,$$

where γ and β are the only two learnable parameters of a batch normalisation layer.

The reason for the scale and shift step is to allow the layer to represent the identity transform if the normalised inputs are not suitable for the following layer, *i.e.* the scale and shift step will reverse the normalisation step if $\gamma = \sqrt{\sigma^2 + \epsilon}$ and $\beta = \mu$. Batch normalisation has now become a standard when training deep CNNs.

3.5.2 Skip-connections

Residual Networks became very popular after it was used to win one of the ImageNet competitions (He *et al.*, 2015a). The residual connection layer can simply be formalised as

$$y = F(x) + x,$$

i.e. combining the input to the layer(s) with the output of the layer(s). Here, the combination is by addition, but other ways can also be used, like multiplication or concatenation. The idea is also illustrated in Figure 3.5.




Figure 3.5: Diagram conceptualising a skip-connection.

The DenseNet (Huang *et al.*, 2016) is another well-known network making use of skip connections. In a DenseNet each layer is connected to every other layer in the network. There are multiple benefits to using skip-connections: they alleviate the vanishing-gradient problem, strengthen feature propagation, encourage feature reuse, and may reduce the number of parameters required (Huang *et al.*, 2016). These benefits show in the empirical evaluations. It is interesting to note that one can draw a parallel between ResNets and boosting methods since both are approaches to fitting a model to the residual (Huang *et al.*, 2017).

3.5.3 Embeddings

An embedding is a layer that maps a discrete input to a numeric vector representation. It was first used in NLP in order to represent words as numbers so that it can be processed by a numeric model. For instance the word “woman”

may be represented by the vector [1, 3, 5] and the word “man” by [2, 4, 6]. The goal is to map discrete inputs to a meaningful vector space where items with similar meaning exist close to each other. In contrast to using a one-hot encoded representation of words, where all words equally far apart. See Figure 3.6 for an illustration of a such space¹.




Figure 3.6: Learned word embeddings in a 2-dimensional space.

Initially the mappings were configured independently of the neural network with approaches based on co-occurrences (Mikolov *et al.*, 2013). The real breakthrough came when the mappings were defined as learnable layers in the network (as done in all of the latest NLP research (Howard and Ruder, 2018, Devlin *et al.* (2018))). Thus they can be tuned just like any other parameter in the network. The parameters of the embedding function (or layer) are first randomly initialised and then gets tuned along with the rest of the neural network during training.

An embedding operation can either be viewed as a table lookup or a matrix multiplication of the discrete input in a one-hot encoded form, *i.e.* $\mathbf{e} = W\mathbf{x}$ where \mathbf{x} is a discrete input in one-hot encoded form and $W : k \times p$ is the matrix containing the embedding where p is the number discrete categories and k is the embedding size. The embedding layer can be reused by all input features with the same input type, making it more efficient and reducing the memory footprint of the model.

¹<https://www.shanelynn.ie/get-busy-with-word-embeddings-introduction/>

3.5.4 Attention

Attention is one of the standout breakthroughs made in deep learning in recent times; especially playing an integral part to NLP's success and other sequence related tasks (Vaswani *et al.*, 2017, Devlin *et al.* (2018)). It was first popularised in neural machine translation (Bahdanau *et al.*, 2014) and now almost used ubiquitously in natural language processing applications. It was also found useful in computer vision applications, like image captioning (Xu *et al.*, 2015), and in audio processing (Duong *et al.*, 2016).

The main idea of an attention module is to force a layer to only focus on a certain subset of its inputs at different stages of computation. For example, in image captioning, one may use a RNN to sequentially output words describing the image. With the use of an attention model, the network is restricted to only look at certain parts of an image at every step, avoiding having to look at the full image everytime. This is illustrated in Figure 3.7. Notice how the network focusses on the bird part of the image when predicting the word "bird" and "flying" and then focussing on the water part of the image when predicting words like "body" and "water". Similarly, when used in machine translation, we can visualise the attention weights to see how the network focusses on a different subset of words for predicting each word in the target language (Figure 3.8).




Figure 3.7: Attention applied to image captioning.

In an ideal world, an attention module should not be necessary and we would want the network to learn if certain parts of images are not required to look at. However, currently this built-in prior proves to be extremely helpful with current learning algorithms.

The core of an attention module can be summarised with these equations:

$$\mathbf{z} = f(\mathbf{x})$$

$$\boldsymbol{\alpha} = \text{softmax}(f_{att}(\mathbf{x}))$$

$$\mathbf{y} = \mathbf{z} \otimes \boldsymbol{\alpha}$$




Figure 3.8: Attention applied to machine translation.

\mathbf{z} are the ordinary activations produced by a layer f given input \mathbf{x} . $\boldsymbol{\alpha}$ are the weights produced by the attention layer, f_{att} after a logit transform so that the weights sum to one. The output of the attention module, \mathbf{y} , is then an elementwise multiplication, \otimes , of \mathbf{z} and $\boldsymbol{\alpha}$, however other combinatorial operations can also be used. More detail can be found in (Xu *et al.*, 2015). We will elaborate on the concepts of *Self-Attention* (Cheng *et al.*, 2016b) and *Multi-Head Attention* (Vaswani *et al.*, 2017) in Section 4.3.1.

3.6 Super-Convergence

The process of setting the hyperparameters for learning algorithms are hard. It requires expertise, extensive trial and error and is more of an art form than a

science. There are many hyperparameters to consider: learning rate, batch size, momentum, and weight decay, to name a few. Although the right parameters make a huge difference in training time and performance, there are no simple and easy ways to find them. Doing a grid search or a random search in the parameter space is computationally very expensive.

In this section we summarise the work done in (Smith, 2018) to give guidance on finding and setting the learning rate, batch size and weight decay. The approach suggested is based on finding the right balance between overfitting and underfitting by using the right hyperparameters. By looking at the validation loss at early stages of training, one will find clues to which parameters are optimal. This eliminates the necessity of running complete grid or random searches.

By following the hyperparameter selection approach discussed next, one can achieve tremendous improvements in training time and accuracy. If done correctly one may exploit a property called *super-convergence* (Smith and Topin, 2017), where neural networks can be trained an order of magnitude faster than with standard training methods. Figure 3.9 plots the test loss over each training iteration of the proposed hyperparameter strategy compared to a fixed learning rate policy. We see that with the proposed method, the model achieves a better accuracy than the standard approach and in eighthth of the training iterations. Next we discuss how to select hyperparameters to achieve super-convergence.

Firstly note that the effects of the learning rate, batch size and other regularisation techniques, like weight decay, are tightly coupled and should be tuned together. We start with setting the learning rate. In Chapter 2 we showed that the learning rate is an important parameter and that one can either keep it fixed during training or decrease it by some factor when it looks necessary. (Smith, 2015) suggests using *cyclical learning rates* (CLR). A CLR cycle consists of two stages, one where the learning rate linearly increases during training up until a certain maximum and one where the learning rate decreases linearly back to a chosen minimum. One needs to specify the number of epochs in each stage, also known as the step size and the minimum and maximum learning rate. One reason why CLR is believed to work is that variable learning rate allows the model to better dodge saddle points on the loss surface.

In his most reason work, (Smith, 2018), the author suggest using only a single cycle of CLR, called the *1cycle* policy. The learning rate setting of the




Figure 3.9: Reduced training iterations and improved performance with super-convergence principle.

1cycle policy is illustrated in Figure 3.10.




Figure 3.10: The learning rate schedule of the 1Cycle policy.

Clearly one also needs to set the bounds of the learning rate, however luckily, the author designed a learning rate range test to help one do just that. In the learning rate range test, training starts with a small learning rate which is

slowly increased linearly throughout a pre-training run. Typically the training loss will slowly start to decrease as the learning rate increases, up until a certain point at which the learning rate becomes too large and the loss starts to increase. The learning rate at this point is the largest learning rate one can use during training. The lower bound of the learning rate can be set to a factor of 10 times less than the max. See Figure 3.11 for example output of the range test and how to determine the learning rate bounds for the 1Cycle policy. One should be careful of too small a stepsize since that increases the rate of the learning rate increase which might make the training process unstable.




Figure 3.11: An example output of a learning rate range test.

Note that large learning rates also act as a form of regularisation (Smith, 2015). Thus if one uses large learning rate, one might want to reduce some of the other regularisation controls. Along with the learning rate (Smith, 2018) suggests and shows that larger batch sizes allow for training with larger learning rates and thus convergence can be reached quicker. So when setting the batch size one would choose the largest possible that fits into memory.

The weight decay is another important value to be set properly along with the 1Cycle policy. Since it is a regularisation parameter it is important that it is in balance with the learning rate. The proposed way of setting the weight decay is through a grid search. Once can do a few short training runs at different weight decays and see if the validation loss gives any indication of which one is better. Another option is to test different weight decays during the learning rate range test and see how they behave. Then we will look for the weight decay that produces the most stable loss and allows the use of the largest learning rate. In our experiments we follow these suggestions to find the best hyperparameters for each model and dataset.

3.7 Model Interpretation

Although Deep Learning is now the state-of-the-art for many machine learning tasks, it is still trailing behind other algorithms in terms of model interpretability. But keep in mind this is not an unusual trade-off; between prediction performance and model interpretability. Deep neural networks are occasionally referred to as “black boxes” since it is very difficult to interpret what is going on inside the stacks of linear and non-linear layers. This is one of deep learning’s greatest criticisms and is a large reason why it cannot be used in some production environments. For example, in the clinical domain, model transparency is of utmost importance, given that predictions might be used to affect real-world medical decision-making and patient treatments (Shickel *et al.*, 2017). Fortunately, some work has been done to gain insights from neural networks. We discuss the topic briefly in the following sections.

3.7.1 Neural Network Specific

We have showed in §2.4 that it is possible to inspect activations and weights of layers at different levels of the network. If the network is small, one might gain insight to what the network has learned or why it is making certain decisions. However, most useful neural networks are at least three layers deep, making its activations and weights more complex to interpret.

When fully convolutional networks are used, there are ways to visualise which parts of the inputs were important in making a certain decision. These visualisations are called class activation maps (Zhou *et al.*, 2016) ((Selvaraju

et al., 2017) suggests a similar approach). But they cannot be used with fully connected layers and only with fully convolutional networks.

Another common interpretation tool, in order to gain insight into what specific neurons are looking for, is to rank the inputs by the magnitude of their activation at that neuron. Then, if one would be able to manually spot similarities between the highest ranked inputs, one would have a potential description of the pattern that triggers that neuron. This approach is called *Activation Maximisation* (Erhan *et al.*, 2009). This process is manual and not guaranteed to produce useful insights, especially when there are many neurons to investigate.

An interesting take on model interpretation is shown in (Frosst and Hinton, 2017) and (Che *et al.*, 2016). Their ideas are based on the process of *knowledge distillation* (Hinton *et al.*, 2015) and leverages the fact that decision trees are easier to interpret. Knowledge distillation is the process of transferring the knowledge from one model (or an ensemble of models) to another model by training the target model to estimate the predictions of the source model (or ensemble). (Frosst and Hinton, 2017) and (Che *et al.*, 2016) used a soft-decision tree and boosted trees respectively to learn the mapping between the input and the neural network predictions. Once the tree based methods are trained, the usual interpretation tools of tree based models, like feature importance of evaluating how a sample traverses the tree, can be used to attempt to understand what the neural network is doing. It is interesting that in both publications the authors note that the tree based models trained on the neural network predictions achieved a better performance than the ones trained on the actual targets - indicating the value knowledge distillation can add.

3.7.2 Model Agnostic

Besides interpretation tools specifically designed for neural networks, one can also make use of model agnostic tools, *i.e.* ones that can be used with any machine learning model. *The Permutation importance algorithm* computes feature importance scores for each input feature. The importance measure of a feature is determined by the model's sensitivity to a random permutation of the values of that feature. The expectation is that when a feature with a strong signal is shuffled before passed to the model, the performance of the model will drop significantly. On the other hand, if a feature has little effect

on the target predictions, shuffling its values will not have a major impact on the model's performance.

This technique was first introduced by (Breiman, 2001) for the Random Forest algorithm but can also be generalised to other models. The steps for calculating the feature importances by the permutation importance algorithm are as follows:

1. Train a neural network on a dataset with p input features.
2. Evaluate the network on a validation set to obtain a performance metric m_0 .
3. For each of the p input features, do:
 - create a copy of the validation set and randomly shuffle the feature in this copy.
 - Evaluate the neural network on this version of the validation set to obtain m_j , $j = 1, \dots, p$.
4. Rank the features based on $m_j - m_0$ (if a bigger m is better).

Permutation importance only produces sensible results on the assumption that the features are independent. Permuting the feature independently creates examples that never occur in real life and the importance of features in that invalid space may be misleading. The test however can still be useful to identify inputs that are not important, *i.e.* not used by the model. If randomly permuting a feature does not affect the model performance at all, it may be a good indication that the model is not dependent on it.

Partial Dependence Plots can be used to visualise the relationship between input features and the target. This is another post-training interpretation tool and was first used in (Friedman, 2001). Once a neural network is trained, we can evaluate what effect a change in any input feature has on a single prediction, by observing the change in the prediction. Say for example we want to see how feature X_1 influences the model prediction. By taking a single observation from the data we can evaluate how the model prediction changes by changing the value of X_1 to other possible values of X_1 . And since this behaviour will most likely vary for different observations, this process should be repeated for a subset of observations from the dataset. The average effect on the predictions at different values of X_1 can then be calculated along with standard errors of these effects.

Other model agnostic interpretation tools, like SHAP values (Lundberg and Lee, 2017), are also available. Here we only discussed the basic approaches. In the next two chapters we will show examples of how these model interpretation techniques can be implemented.

Chapter 4

Deep Learning for Tabular Data

4.1 Introduction

Up to now we have covered the basics of neural networks, as well as the more recent advancements in the era of deep learning. The aim of this chapter is to explore how the approaches and techniques we discussed can be leveraged when doing deep learning with tabular data. Tabular data is vastly different to so-called unstructured data like images, text and speech and therefore we do not expect all of the methods to be as effective. Deep learning on tabular data is not as well researched compared to the aforementioned data domains and thus it is not always clear how to solve certain modelling challenges. There is only a handful of publications reporting successful implementations of deep learning for tabular data on important applications such as recommender systems (Haldar *et al.*, 2018), click-through rate prediction¹ (Song *et al.*, 2018), analysis of electronic health records (Rajkomar *et al.*, 2018) and transport related problems (de Brébisson *et al.*, 2015), to name a few.

The tabular data domain is still dominated by tree-based models like random forests and gradient boosted trees. This makes us curious as to why deep learning is not as effective here as it is in most other data domains. This chapter will help shed light on this issue and point to promising ways of improving the current state-of-the-art.

This chapter is structured based on the tabular dataset challenges introduced

¹To predict the probability of a user clicking an item, critical to online applications.

in §1.3. For each of the challenges we provide a recap of the issue, review the literature to discuss how they are currently being treated, and provide suggestions for improving these approaches, (where applicable). In §4.2 we look at ways of representing the input features of a tabular dataset. §4.3 is about approaches for leveraging feature interactions. A large part of this chapter is on how to be more sample efficient; this we discuss in detail in §4.4. Then we briefly investigate ways of interpreting deep neural networks for tabular data in §4.5. In the final section §4.6 we discuss the 1Cycle policy and hyperparameter selection for tabular data neural networks, in addition to other miscellaneous topics that do not fit into the above categorisation.

4.2 Input Representation

One of the major design considerations when building a deep neural network for tabular data is the input representation, *i.e.* how should one represent each feature numerically? This choice may heavily influence the model’s ability to extract patterns from the input and the effectiveness of the optimisation algorithm. What makes this decision harder is that the features in a tabular dataset are highly heterogeneous (Shavitt and Segal, 2018). What we decide for one features might not be optimal for another and we want to ensure that no feature dominates another in the training process.

A tabular dataset typically has continuous features and categorical features. Different approaches are needed to process each. But the processing should be done in such a way that the one feature type does not dominate the other in terms of feature extraction and optimisation.

A tabular dataset can be high-dimensional and very sparse which makes the task more difficult but even more important to do effectively, as noted by many (Song *et al.*, 2018, Wang *et al.* (2017), Qu *et al.* (2016), Cheng *et al.* (2016a), Anonymous (2019), Covington *et al.* (2016)). An example of one such challenging tabular dataset is the Criteo dataset². Its feature dimension is ~ 30 million, exhibiting sparsity of $\sim 99\%$. Although not always as extreme, these properties are common in tabular datasets.

²<https://www.kaggle.com/c/criteo-display-ad-challenge>

4.2.1 Numerical Features

One of the things that make tree-based methods so attractive is that the scale and the distribution of the features hardly matter, as long as their relative ordering is meaningful. With neural networks, we are not that fortunate. Neural networks are very sensitive to the scale and distribution of its inputs (Ioffe and Szegedy, 2015). If the features are on different scales, one might dominate the weight updates, and if a feature contains a large value, it may throw off the optimisation procedure and cause exploding or vanishing gradients (Clevert *et al.*, 2015). This means that one should do a proper standardisation of all continuous features in a tabular dataset.

The typical standardisation approach for numerical features in deep learning is to do mean centering and variance scaling, *i.e.* $\tilde{x} = (x - \mu)/\sigma$, where μ and σ is the mean and standard deviation of X respectively ($x \in X$). One would expect this transformation to be sufficient, but in practice, many have found otherwise.

In (Haldar *et al.*, 2018) they suggest to first inspect the distribution of each of the features. If a feature looks gaussian, do the standard normalisation, $(x - \mu)/\sigma$, but if the feature looks more like a power law distribution, transform it by $\log((1 + x)/(1 + \text{median}))$. This ensures that the bulk of the values lie between $\{-1, 1\}$ and that the median is close to zero. As an example, we illustrate what effect these transformations have on 2 continuous features in the Adult dataset, Age and Hours-per-week, in Figure 4.1. In Figure 4.1 (a) we see that the features are on a roughly similar scale but their distributions are totally different. By applying mean centering and unit variance scaling (b) we see the values of the Hours-per-week feature are mostly in $\{-1, 1\}$ but for the Age features there are still plenty of values outside that range. When doing the power distribution transformation of (Haldar *et al.*, 2018), we observe that the Age feature is now in the proposed value range. The downside to this approach is that it is not automatic and can become cumbersome with many variables.

To reduce the possibly high variance of numeric features (Song *et al.*, 2018) suggests to transform the a numeric feature to $\log^2(x)$ if $x > 2$. This worked in their use-case but it is hard to imagine that this work on all other cases. This transformation causes a discontinuity at $x = 2$ and a possible overlap between values the were originally less than 2 with those that were greater than 2. In addition, this transformation does not address the extreme values on the negative side. (Wang *et al.*, 2017) simply used a standard log transform,




Figure 4.1: Effect of normalisation on continuous variables.

$\log x$, to normalise continuous features.

In (Covington *et al.*, 2016) they also found that the proper normalisation of numeric features were critical for the model to converge. Their approach was to transform the numeric features to be equally distributed in [0,1) using the cumulative distribution $\tilde{x} = \int_{\infty}^x df$, where f is the distribution of x . They approximated the integral with linear interpolation on the quantiles of the feature values computed as a preprocessing step.

Another possibility for numeric feature normalisation that we have not seen used in the deep learning for tabular data literature, is to use a batch normalisation layer as the initial layer to numeric features. This will apply the same type of scaling as the zero mean and unit variance transformation, but with transformation parameters learned from each batch. Thus the need to preprocess numeric features is removed. The caveat is that the quality of the transformation depends on the batch statistics and thus if the batch is not representative of the full data distribution (which is likely if the batch size is small), then it might derail the training procedure.

The fact is that the only way to know for sure which normalisation to apply to numeric features is by experimentations, since it varies between datasets. The standard normal scaling is a safe transformation to use if experimentation wants to be avoided.

4.2.2 Categorical Features

Most of the sparsity in tabular datasets come from categorical features. Since neural networks cannot process discrete categories or objects, we need to find

a numeric representation for each class. The standard approach is to one-hot encode categorical features. That is, if we have a categorical feature with three possible categories, for example, the one-hot encoded form of the three categories will be: Category 1: [1, 0, 0], Category 2: [0, 1, 0] and [0, 0, 1] for Category 3.

There are multiple inefficiencies when using one-hot encodings with neural networks. The obvious one we have already mentioned is that it introduces sparsity to the data, since the dimension of the one-hot encoded form is equal to the number of categories in a feature. Thus if we have plenty of high-cardinality features in our dataset, the data will be extremely sparse and difficult to model. This also increases the size needed of the first linear layer and thus we need machines with greater memory and processing power. Neural networks can easily overfit such sparse data (Covington *et al.*, 2016).

The other problem with one-hot encodings of categorical features is that there is no notion of similarity and distances between categories. In this representation, all categories are equally far apart now matter how semantically similar or dissimilar they are. This makes it harder for the model to learn useful patterns.

An alternative to one-hot encodings as representations of categorical features for neural networks are *entity embeddings*. An entity embedding is the exact same operation as (word) embeddings we have discussed in §3.5.3, but applied to categories instead of words. Therefore an entity embedding assigns a numeric vector representation to each category in a categorical feature, for example: Category 1: [0.05, -0.1, 0.2], Category 2: [0.2, 0.01, 0.3] and [-0.1, -0.2, 0.05] for Category 3. Once all of the categorical features have been embedded, their representations can be concatenated and passed to the rest of the network.

The first published work in modern times on entity embeddings was in the taxi destination prediction challenge (de Brébisson *et al.*, 2015). In another Kaggle success story, (Guo and Berkhahn, 2016) successfully used entity embeddings for predicting the total sales of a store. Companies like Instacart and Pinterest have reported the effective use of entity embeddings on their internal datasets. Currently, all research on deep learning for tabular data makes use of entity embeddings - see for example (Song *et al.*, 2018), (Wang *et al.*, 2017), (Covington *et al.*, 2016) and (Zhou *et al.*, 2017).

The reason why it is used all-round is because it does not have the same issues as the one-hot encoded representations. A similar formulation to §3.5.3,

we define the embedding for the j -th categorical feature by:

$$\mathbf{e}_j = V_j \mathbf{x}_j$$

where \mathbf{x}_j is the one-hot encoded vector representation of the j -th categorical variable and V_j is the associated embedding/weight matrix. The weights in V_j are learned along with all of the other parameters in the network.

The first advantage of entity embeddings is that it speeds up training and reduces the memory footprint which further improves the generalisation ability of the network (Covington *et al.*, 2016), (Guo and Berkhahn, 2016). This especially useful when working with dimensional and sparse inputs. Suppose we have a dataset with two categorical features, X_1 and X_2 , with cardinality of C_1 and C_2 , respectively. Furthermore, suppose that the first hidden layer in the neural network accepts inputs of size q and thus we need project an observation with these two features into a vector representation of the same size. If we were to use one-hot encoded representations of X_1 and X_2 we would need a weight matrix of size $(C_1 + C_2) \times q$. However, if we use entity embeddings of X_1 and X_2 we may have two weight matrices of sizes $C_1 \times q/2$ and $C_2 \times q/2$, which is in total half the number of parameters needed compared to the pure one-hot encoded representation.

The size of the embeddings is a hyperparameter of the model and again there is no way to tell what this value should be beforehand. Most publications rely on a grid search to find the optimal sizes. For example (Song *et al.*, 2018) experimented with embedding sizes [8,16,24,32] and found 16 to work the best, and (Cheng *et al.*, 2016a) found that an embedding size of 32 was optimal for their use-case. The selection complete depends on the data and the network used.

(Wang *et al.*, 2017) and (de Brébisson *et al.*, 2015) used different embedding sizes for each categorical features and suggested these rules-of-thumb:

- $6 \times (\text{cardinality})^{\frac{1}{4}}$ (Wang *et al.*, 2017)
- $\max(2, \min(\text{cardinality}/2, 50))$ (de Brébisson *et al.*, 2015)

It makes sense to have different embedding sizes for categorical features with different complexities.

Entity embedding not only reduces memory usage and speeds up neural networks compared to one-hot encoding, but more importantly, by mapping similar values close to each other in the embedding space it reveals the intrinsic

properties of the categorical variables, which one cannot obtain with one-hot encoding. This allows us to interpret the classes of the categorical features. The embeddings can be visualised to gain further insight into the data and model decision making. The weights associated with each category's projection onto the embedding space can be plotted with any dimension reduction technique like t-sne or PCA. Then we can compare the categories based on their relative distances and positions in this reduced space. In Figure 4.2 we plot a 2 component PCA of the embedding matrix of the Education categorical feature in the adult dataset. We see that the school categories all lie in the bottom-right corner of the space, with some notion of ranking from grade 5 (top-right) to grade 12 (bottom-left). The tertiary education classes are in a separate cluster and their levels of education correspond conveniently with the vertical axis.




Figure 4.2: PCA of the Education entity embedding weight matrix.

To prove that these entity embeddings actually learns something useful, besides plotting the embedding matrix, one can also feed them along with the continuous features to other learning algorithms and see how it affects performance. (Guo and Berkahn, 2016) found that the embeddings obtained from the

trained neural network boosted the performance of all tested machine learning methods considerably when used as the input features. These embeddings can be reused on different machine learning tasks and do not have to be re-learned for each dataset. Instacart and Pinterest, referenced above, reported sucessful implementation of this approach.

The entity embedding approach is very flexible. One can reuse an embedding over different categorical features if the features have overlapping categories. (Zhou *et al.*, 2017) has an intersting take on multi-hot categorical features; where a feature can have more than category associated with it. The embedding layer for that instance then outputs a list of embeddings with length the same as the number of categories associated with that instance and feature. The list of embeddings then gets projected back into a fixed-length representation by doing a pooling operation.

4.2.3 Combining Features

Once the continuous and categorical features have been processed and embedded, we need a way to combine them before passing it to the rest of the network. The standard approach is to concatenate each categorical variable embedding to the continuous variables, as was done in (Haldar *et al.*, 2018) and (Wang *et al.*, 2017) for example, illustrated in Figure 4.3. The potential problem with this approach is that some features might be over-represented in this vector. For example, one of the continuous features may be very important for prediction, but it gets lost when concatenated with all the entity embeddings each taking up more space in the combined representation.

In (Song *et al.*, 2018) they embed both the numerical and categorical features into the same sized embedding. By mapping both types in the same feature space facilitates more effective learning of interactions between the mixed features. The embedding for the j -th numerical features is obtained by:

$$\mathbf{e}_j = \mathbf{v}_j x_j$$

where x_j is a scalar and \mathbf{v}_j is the associated weight vector.

4.3 Learning Feature Interactions

In most machine learning tasks the greatest performance gains can be achieved by feature engineering wheras better algorithms may only result in incremental




Figure 4.3: Combined representation of continuous and categorical features.

boosts. In feature engineering one strives to create new features from the original features based on some domain knowledge of the data or otherwise, that makes it easier for the model to estimate the target and help capture high-order interactions between features. Although a crucial step to make the most out of the data, this can be a very laborious process. It is widely held that 80% of the effort in an analytic model is preprocessing, merging, customizing, and cleaning datasets, not analysing them for insights (Rajkomar *et al.*, 2018). There is no formal path to follow in this stage and thus usually consists of many a trial and error, benefitted by domain knowledge of the data, which is not always accessible. A huge advantage of using neural networks on tabular data (and other data structures) is that the feature engineering process gets automated to some extent (with the caveat of numerical feature preprocessing.). A neural network learns these optimal feature transformations and interactions implicitly during the training process. The hidden layers of a neural network can be viewed as a feauture extractor that was optimised to map the inputs into the best possible feature space for a model (the final layer of the network) to operate in.

The standard approach is to stack a few fully-connected layers one after the other to map the input representation to the output, as was done in (Covington *et al.*, 2016). Fully-connected layers model all feature interactions implicitly and in ideal world, we would expect this architecture to be sufficient. In practice, with out current learning algorithms, a simple MLP is not good enough to learn all types of interactions. A fully connected model structure leads to very complex optimisation hyperplanes with a high risk of falling into local

optimums.

Therefore it is necessary to explicitly leverage expressive feature combinations or encourage the network to learn better high-order feature interactions. This issue receives attention in publications such as (Song *et al.*, 2018), (Wang *et al.*, 2017), (Qu *et al.*, 2016) and (Guo *et al.*, 2017). The restrictions we impose on the fully-connected structure may further help to limit the model size to make learning more efficient. The key question we asked in this section is how do we help the network to determine which features to combine to form meaningful high-order features. We first briefly review some of the suggestions in the literature, after which we focus on the attention mechanism which we believe is the most powerful method of learning feature interactions.

The authors of (Wang *et al.*, 2017) make a case for finding a bounded-degree of feature interactions, arguing that all the Kaggle competitions are won with feature engineering of low-degree interactions, whereas deep neural networks learn highly non-linear interactions implicitly. They suggest using an automated way of building cross-features, called the cross-network. In the cross-network each layer produces higher-order interactions based on existing ones, and keeps the interactions from previous layers. The cross-network consists of cross-layers that can be formalised as:

$$\mathbf{x}_{l+1} = \mathbf{x}_0 \mathbf{x}_l^\top \mathbf{w}_l + \mathbf{b}_l + \mathbf{x}_l$$

where \mathbf{x}_l is the output of the l -th cross layer and the input to the $(l + 1)$ -th cross-layer; \mathbf{x}_0 is the combined input representation; \mathbf{w}_l and \mathbf{b}_l are its associate weight and bias parameters respectively. Each cross-layer adds back its input after feature in the same fashion as a skip-connection. The degree of the cross-features grows with cross-network depth. The authors experimented with 1-6 cross-layers and found that a depth of 6 gave the best results.

In parallel to the cross-network, they also used a standard deep neural network to learn the highly non-linear feature interactions. The DNN accepts the same input and its output is then concatenated with that of the cross-network. These two networks can then be trained simultaneously.

(Qu *et al.*, 2016) used something called a product layer, which takes pairwise inner or outer products of all feature combinations and concatenates it to all linear combinations. The output is then fed to 2 fully-connected layers.

According to (Guo *et al.*, 2017) and (Cheng *et al.*, 2016a) it is necessary to capture both low and high-order interactions. They both achieve this by

having two parallel networks, similar to (Wang *et al.*, 2017), where one learns high-order interactions and the other low-order interactions. Both use a deep neural network to learn the high-order interactions. (Cheng *et al.*, 2016a) uses a shallow but wide neural network to capture the low order interactions and where (Guo *et al.*, 2017) uses a learnable factorisation machine to do that. Similar to the other parallel stream networks, the output of the two streams gets concatenated before passing it to the classification layer.

4.3.1 Attention

Based on the results in the paper (Song *et al.*, 2018) and by our findings in general deep learning research §3.5.4, we deem attention to be the most promising mechanism to model feature interactions. (Song *et al.*, 2018) uses a multi-head self-attention mechanism which they call the interacting layer. Within in the interacting layer each feature is allowed to interact with every other feature and automatically determine which of those interactions are relevant to the output.

To explain the attention mechanism, consider feature j and suppose we want to determine which high-order features involving feature j are meaningful. We first define the correlation between features j and k under attention head h as:

$$\alpha_{j,k}^{(h)} = \frac{\exp(\phi^{(h)}(\mathbf{e}_j, \mathbf{e}_k))}{\sum_{l=1}^L \exp(\phi^{(h)}(\mathbf{e}_j, \mathbf{e}_l))}$$

where $\mathbf{e}_l, l = 1, \dots, L$ is the embedding of the l -th features. $\phi^{(h)}(\cdot, \cdot)$ is an attention function which defines the similarity between two features. It can be defined by a trainable layer or a simple inner product as in:

$$\phi^{(h)}(\mathbf{e}_j, \mathbf{e}_k) = \langle W_{\text{query}}^{(h)} \mathbf{e}_j, W_{\text{key}}^{(h)} \mathbf{e}_k \rangle$$

where $W_{\text{query}}^{(h)}$ and $W_{\text{key}}^{(h)}$ are transformation matrices which map the original embedding space into new spaces of the same dimension. The representation of feature j in subspace h is then updated by combining all relevant features guided by the coefficients $\alpha_{j,k}^{(h)}$:

$$\tilde{\mathbf{e}}_j^{(h)} = \sum_{k=1}^K \alpha_{j,k}^{(h)} W_{\text{value}}^{(h)} \mathbf{e}_k$$

$\tilde{\mathbf{e}}_j^{(h)}$ is a combination of feature j and its relevant features under attention head h . Note that $\alpha_{j,k}^{(h)}$ for $k = 1, \dots, K$ sum to 1 since they went to a logit transform/softmax operation.

$\tilde{\mathbf{e}}_j^{(h)}$ is thus a learned combinatorial feature. Since a feature can be involved in various different combinations, we use multiple heads to extract the multiple combinations, *i.e.* $\{\tilde{\mathbf{e}}_j^{(h)}\}_{h=1}^H$. In the original paper (Song *et al.*, 2018) used $H = 2$, but this is typically a hyperparameter one needs to tune. All of these combinatorial features are concatenated into a single vector, $\tilde{\mathbf{e}}_j$. Then finally the output is combined with its raw input and sent through a ReLU activation:

$$\mathbf{e}_j^{\text{res}} = \text{ReLU}(\tilde{\mathbf{e}}_j + W_{\text{res}}\mathbf{e}_j)$$

This mapping from \mathbf{e}_j to $\mathbf{e}_j^{\text{res}}$ is done for each feature to form the interacting layer. The activations of the interacting layer is thus a representation of the high-order features of its inputs. These interacting layers can be stacked on-top of each other to form arbitrary order combinatorial features.

(Zhou *et al.*, 2017) follows a vaguely similar idea to attention by learning the weights to apply to the hidden representations. Interestingly, they removed the softmax layer as a way to mimick probabilities to reserve the intensity of activations.

In terms of the skip-connection already mentioned a couple of times above - both (Song *et al.*, 2018) and (Wang *et al.*, 2017) found by experimentation that by having skip-connections to connect lower-level features with higher-level ones improves the performance of the network.

4.3.2 Self-Normalising Neural Networks

By studying the literature of deep neural networks for tabular data we rarely see the optimal network depth go beyond three or four layers. The reason is that a fully connected model have very complex optimisation hyperplanes which increases the risk of falling into bad local optima. A proposed way of training deeper neural networks is to make use of *Self-Normalising Neural Networks* (Klambauer *et al.*, 2017). They were developed as an alternative to batchnorm layers since batchnorm layers often become unstable when using SGD or stochastic regularisation techniques like dropout, especially when using fully-connected neural networks. This is exhibited by a high-variance in the training error.

The self-normalising neural network is simply a neural network with a novel activation function called the *SeLU*. The SeLU helps the network to maintain zero mean and unit variance for the activations at all network levels. The SeLU activation eliminates the need for a BatchNorm layer and is also much safer against exploding or vanishing gradients. It is defined as:

$$\text{selu}(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha e^x - \alpha & \text{if } x \leq 0 \end{cases}$$

where $\lambda = 1.0507$ and $\alpha = 1.6733$ are special constants derived in the paper. When using the SeLU activation function, one should be aware of the special weight initialisation and dropout technique that should be used in tandem.

(Klambauer *et al.*, 2017) tested SeLUs on 121 classification datasets from the UCI Machine Learning repository. They compared DNNs with SeLU activations to other DNNs and other classifiers like Random Forests and SVMs. They found that on the datasets with less than 1000 observations, random forests and SVMs performed the best. However, for the datasets with more than 1000 observations, DNNs with SeLU activations performed the best overall. The classifiers were compared by ranking them by their accuracy for each prediction task and doing a pairwise Wilcoxon test. Another thing the authors found when comparing SeLUs with other activations is that the model selection approach for DNNs with SeLUs resulted in much deeper networks than DNNs with other activation functions.

These are all promising results but we have not seen SeLUs been used in other applications or papers on tabular datasets.

4.4 Sample Efficiency

It is well known that deep neural networks require a large amount of data to generalise well. Typically, tabular datasets are not as large as unstructured datasets like images and texts. There is also no large tabular dataset from which knowledge can be transferred, like an ImageNet for computer vision or a Wikipedia corpora for NLP equivalent. We suggest two techniques for overcoming this problem: data augmentation and unsupervised pre-training.

4.4.1 Data Augmentation

Data augmentation for tabular datasets is rarely studied. Tabular data is very different to image data and the standard augmentations used in computer vision does not make sense with tabular data. You cannot rotate or scale an observation from a tabular data without losing its meaning. One transformation that does make sense for tabular input is the injection of random noise.

When working with images, we can randomly perturb the pixel intensities by a small amount so that it is still possible to make sense of its content. By adding 1 for example to all pixels and all colors in an image, will only make it slightly brighter and we will still be able to make sense of it. But with tabular data we can just randomly add a small amount to any feature. The input features will probably not all be on the same scale and the addition of noise might result in a feature value that is not within the true distribution. In addition, it does not make sense to add anything to a discrete variables. Thus in order to inject random noise to a tabular data sample, the noise should be scaled relative to each input feature range and the results should be a valid value for that feature. This also helps the model to be more robust to small variations in the data. (Van Der Maaten *et al.*, 2013) suggests an augmentation approach that does this called *Marginalised Corrupted Features* (MCF). The MCF approach adds noise to input from some known distribution. The process is manual and it cannot be applied to discrete variables.

In the original DAE paper (Vincent *et al.*, 2008), they used a blank-out corruption procedure. Which is randomly selecting a subset of the input features and masking their values with a zero. This similar to dropout regularisation but instead applied to the inputs. The only conceptual problem with this approach is that for some features a zero value actually carries some meaning, so a suggestion is to blank-out features with a unique value not already belonging to that feature distribution.

Another input corruption approach shown to work empirically here is what is a technique called *Swap Noise* (Kosar and Scott, 2018). The swap noise procedure corrupts inputs by randomly swapping input values with those of other samples in the datasets. We have illustrated the approach in Table 4.1. In this way you ensure that the corrupted input have at least valid feature values. But it still might produce combinations of features that are not actually possible according to the true data distribution.

The final data augmentation we look at is *mixup augmentation* (Zhang *et al.*,

Table 4.1: Swap Noise Example.

age	occupation	education	race	sex	$\geq 50k$
Original Dataset					
49	NA	Assoc-acdm	White	Female	1
44	Exec-managerial	Masters	White	Male	1
38	NA	HS-grad	Black	Female	0
38	Prof-specialty	Prof-school	Asian-Pac-Islander	Male	1
42	Other-service	7th-8th	Black	Female	0
20	Handlers-cleaners	HS-grad	White	Male	0
Sample with swap noise					
49	Prof-specialty	Prof-school	Asian-Pac-Islander	Female	1

2017). The way mixup creates artificial samples is by the following formulation:

$$\tilde{\mathbf{x}} = \lambda \mathbf{x}_i + (1 - \lambda) \mathbf{x}_j \quad \tilde{\mathbf{y}} = \lambda \mathbf{y}_i + (1 - \lambda) \mathbf{y}_j$$

where \mathbf{x} is a input vector, \mathbf{y} a one-hot encoded output vector and $\lambda \in [0, 1]$. $(\mathbf{x}_i, \mathbf{y}_i)$ and $(\mathbf{x}_j, \mathbf{y}_j)$ are two samples drawn at random from the training data. Thus mixup assumes that linear interpolations of input vectors lead to linear interpolations of corresponding targets. We visualise the creation of artificial samples through mixup on a toy dataset in Figure 4.4.




Figure 4.4: Illustration of points created by mixup augmentation.

λ controls the strength of the interpolation between the input-output pairs. The closer λ is to 0 or 1, the closer the artificial sample will be to an actual training sample. The authors suggest using $\lambda \sim \text{Beta}(\alpha, \alpha)$ for $\alpha \in (0, \infty)$, since this makes it more likely that the mixed-up sample lie closer to one of the original samples than in the middle of the two. They observed best performance when $\alpha \in [0.1, 0.4]$ and if α is too high, they experience underfitting.

Other ablation studies they did was to find at which stages of the network the interpolation should happen, *e.g.* raw input, after embedding, hidden layer, *etc.* But the experiments are not extremely clear and therefore warrants further discussion here.

Typically data augmentation procedures are dataset dependent and therefore requires expert knowledge. It is hard to think of ways to augment tabular data, even more so a generic way of doing so. However, from this definition it is clear that mixup can be used on any type of data, including tabular datasets.

Mixup data augmentation can be understood as a mechanism to encourage the model to behave linearly in-between training samples. (Zhang *et al.*, 2017) shows that this linear behaviour reduces the amount of undesirable variation when predicting new samples further away from the training samples. They also argue and show empirically how training with mixup is more stable in terms of model predictions and gradient norms. This is because mixup leads to decision boundaries that transition linearly between classes, resulting in smoother predictions.

The authors tested mixup data augmentation on tabular datasets. They tested it on six classification datasets from the UCI Machine Learning repository. They used a 2-layer MLP with 128 neurons each and a batch size of 16. They found that mixup improved the performance on four out of the six datasets.

All of these methods have hyperparameters that needs to be set. For example, with swap noise and blank-out we must set the proportion of feature values to be corrupted and with mixup we need to set the mixup weight. These parameters control the strength of the regularisation and thus should be tuned accordingly.

4.4.2 Unsupervised Pretraining

We described unsupervised pretraining in . It is a very effective way of training deep neural networks. It also addressed the feature engineering issue we face

with tabular data by automatically learning useful feature combinations in an unsupervised fashion. However, we only found two published papers using this approach for tabular data, (Zhang *et al.*, 2016) and (Miotto *et al.*, 2016). (Miotto *et al.*, 2016) used a stacked denoising autoencoder as a feature learning method to derive a general-purpose patient representation for EHR data that facilitates clinical predictive modelling. The trained DAE is used as a fixed feature extractor to obtain representations of the data that can be passed to supervised learning algorithms. (Zhang *et al.*, 2016) also did pretraining with DAEs.

Feature learniattempts to overcome limitations of supervised feature space definition by automatically identifying patterns and dependencies in the data to learn a compact and general representation that make it easier to automatically extract useful information when building predictive models (Miotto *et al.*, 2016). These techniques are very familiar and effective in text, audio and image processing, but not with tabular data.

The DAEs enforce the representations to be robust to partially destroyed inputs. By using these learned representations as input, significantly improved the performance of predictive models compared to those only using the raw inputs.

There are a lot of unknowns when it comes to using DAEs for pretraining. First we need to decide on the structural hyperparameters of the network (number of layers, sizes, etc.) and then we need to decide what learning paramters to choose. This is not as thoroughly explored as hyperparamter selection for supervised neural networks. We also need to decide which loss function we need to used for the DAE. This choice depends on the type of the input. If all the inputs are numeric we can use the common MSE loss, but if the input has categorical features, this decision needs more consideration. (Miotto *et al.*, 2016) transformed all numeric features to be in the range of [0,1] and one-hot encoded the categorical features so that they can use the binary cross-entropy as the loss function. But we cannot follow this approach if we wish to use entity embeddings for the categorical features. A suggestion is to have a separate multiclass cross-entropy loss for each categorical feature and use the MSE loss for the continuous features. These losssed should then be combined to obtain a total loss. How to weight the different losses is also something that is not yet researched.

We suggest treating the unsupervised pretraining rather as a transfer learn-

ing problem than a feature extraction method like in (Zhang *et al.*, 2016) and (Miotto *et al.*, 2016). That is, slicing the head of the trained DAE and replacing it with the required classification layers. Then first train only the new layers keeping the lower level DAE layers fixed and then train all of the layers simultaneously. By doing it this way, the “feature extractor” part of the DAE can still be tweaked to be optimal for the supervised learning task. We experiment with this approach in the next chapter.

4.4.3 Regularisation

Since regularisation can also help one to be more sample efficient we have a brief discussion about it here. A popular form of regularisation in the papers we have studied is dropout (Song *et al.*, 2018), (Zhang *et al.*, 2016), (Qu *et al.*, 2016), (Guo *et al.*, 2017). The dropout percentage was mostly determined using a grid search and the results showed the optimal percentage depends on the data and the model being used.

Interesetingly, (Haldar *et al.*, 2018) found that dropout was not effective in their application. They pinned it down to dropout producing invalid input scenarious that distracted the model. Therefore they opted for hand crafted noise shapes taking into account the distribution of the relevant feature.

The other form of regularisation used for tabular data is weight decay (Song *et al.*, 2018), (Wang *et al.*, 2017), (Zhang *et al.*, 2016), (Qu *et al.*, 2016) and (Zhou *et al.*, 2017), (Zhang *et al.*, 2016) compared dropout with L2 and found dropout to be better.

For a different approach to regularisation, recall the difference between tabular data and unstructured data. A difference between the two data types that stand out is the relative importance of each of the important features with respect to the target. In computer vision a large amount of pixels should change before an image is of something else. Whereas in tabular data a very small change in a single feature may have totally different behaviour with respect to the target (Shavitt and Segal, 2018). The same authors mention that this can be addressed by including a separate regularisation term for each of the weights in the network. These regularisation terms are seen as additional model hyperparameters. It is easy to see that this approach is totally intractable since the only way to train these hyperparameters are brute force and repetitive tweaking and validation which cannot be done in

by a gradient based method. A workaround is to make these regularisation parameters trainable like all of the other points in the network. This is achieved by minimising the counterfactual loss, a novel loss function proposed by (Shavitt and Segal, 2018). They found that training neural networks by optimising the counterfactual loss, outperform other regularisation approaches for neural networks and results in neural networkss that are comparable in performance to gradient boosted trees.

4.5 Interpretation

We have mentioned that model explainability is important for various reasons. It helps one to know how to improve one's model or where it goes wrong. One can even go so far as to say if you cannot explain how a prediction is made, you cannot know how good it is. In the literature we focussed on for this work, not much attention was given to model interpretability. Although there are a few extra interpretation options available when using some of the abovementioned approaches to tabular data.

Firstly, if one uses entity embeddings, the embeddings can be visualised once projected into a low dimensional space, as done in Figure 4.2. (Zhou *et al.*, 2017) is an example in the literature that did this. The attention layer can also provide insight into which feature combinations are important. (Song *et al.*, 2018) and (Zhou *et al.*, 2017) made use of this information. The information is obtained from the attentional weights α which is based on feature similarity and can indicate which features are combined to make predictions for certain observations. Example output in (Song *et al.*, 2018) from this approach can be found in Figure 4.5. Also see (Shavitt and Segal, 2018) for how the learned regularisation parameters were used to evaluate feature importance.

(Haldar *et al.*, 2018) used one of the model agnostic model interpretation algorithms. They used the permutation importance algorithm to evaluate the importance of features. However, they came to the same conclusion as what we stated in §3.7.2. That is that the permutation test did not produce sensible results because the features were not independent. Permuting the feature independently created examples that never occurred in real life, and the importance of features in that invalid space misled them. They did find it somewhat useful to identify redundant features.

We implemented an example of a permutation importance plot for a trained




Figure 4.5: The attention weights visualised for a singe observation in the dataset.

neural network on the Adult dataset. We shuffled each feature 5 times to get a more robust estimate and to obtain a standard error and we measured the performance in terms of the log loss. The resulting plot is in Figure 4.6. The standard error are too small to show on the figure. The horizontal axis measures the average drop in the log loss when the feature is shuffled before prediction. The ranking of the features can be justified intuitively. It is reassuring to see that sex, race and native-country are some of the less important features when it comes to predicting the income of an individual.

Next we compare the permutation importance with the feature importance obtained from knowledge distillation to a boosted trees algorithm (see §3.7.1). To recap the process, we use the predictions of the neural network trained on the Adult dataset as the target for a boosting model to predict using the same input as that given to the neural network. Once the boosting model is trained on these targets, we can extract the feature importance based on the




Figure 4.6: Permutation importance plot of a neural network trained on the Adult dataset.

number of times the base trees used a certain feature to split on. The results are displayed in Figure 4.7.

The ranking of the variables between the two approaches are somewhat the same, except for the marital-status and fnlwgt features. This difference might firstly be caused by the permutation importance method not picking up on multivariate feature interactions and thus incorrectly thinks that marital-status is an important feature, but in reality the performance drop was actually caused by another feature dependent on marital-status. The other explanation we can give is that the decision trees are more likely to split on features with more splitting points and since the fnlwgt feature is the feature with most unique values in the dataset, it contributed to the reason why it was considered so important for the boosting model. These two examples show that one should take care when interpreting models and should use more than one source of information when possible.




Figure 4.7: Feature importance obtained from a boosted trees model trained on the neural network predictions.

4.6 Hyperparameter Selection

We have already identified a plethora of hyperparameters to be tuned if we want to successfully fit a deep neural network to tabular data. Broadly speaking, there are structural hyperparameters to select, like the number of layers and the type of activation function, but then there are also learning algorithm hyperparameters, like the learning rate or the weight decay. In terms of the structural parameters, since there are no shared patterns among the diverse tabular datasets, it is hard to design a universal architecture that will fit all. In most of the literature we have investigated the researcher performs a grid search over many combinations of hyperparameters (Song *et al.*, 2018), (Wang *et al.*, 2017), (Zhang *et al.*, 2016), (Qu *et al.*, 2016), (Guo *et al.*, 2017), (Covington *et al.*, 2016). Most of these parameters are very dependent on the dataset and other modeling choices and therefore the need to tune them. The main structural hyperparameters tuned in the majority of the papers are: hidden layer size, number of hidden layers, activation functions and the shape of the fully-connected layers (*i.e.* is it a constant shape, increasing, decreasing or a diamond shape). The values found varied across publications, showing again

that the step is necessary when working with custom datasets and models. We include some similar experiments in Appendix B

In terms of the learning algorithm hyperparameters, not much experimentation was shown in the publication itself. The only choice that were shown was selecting the learning rate over a grid of values, as in (Zhang *et al.*, 2016) and (Wang *et al.*, 2017), and kept the selection constant during training. Most of the work used the Adam optimiser and early stopping to prevent overfitting (Song *et al.*, 2018), (Wang *et al.*, 2017), (Zhang *et al.*, 2016). Large batch sizes were chosen in (Song *et al.*, 2018), 1024 and (Wang *et al.*, 2017), 512.

Clearly, the 1Cycle policy and the concept of superconvergence has not yet been tested in the tabular data setting. Therefore we will test its effectiveness here using the Adult dataset. We compare it against a constant learning rate. The evaluation method is described in chapter 5. We first do a learning rate range test and find that the optimal learning rate bounds for the one-cycle policy is 0.01 and 0.1. In our experiment we compare training a neural network with a fixed learning rate at the lower bound, 0.01, a neural network with a fixed learning rate at the upper bound, 0.1 and a neural network with a 1Cycle learning rate schedule in these bounds. The results are displayed in Figure 4.8. From the green line, we see that when training with too large of a (constant) learning rate, the validation loss struggles to converge. Training with the smaller constant learning rate (blue line) works better in this case, but the losses plateau quite early on and shows no sign of improving. The 1Cycle learning rate update policy (orange line) shows the best validation loss and accuracy of the three methods. This proves that the 1Cycle policy is also effective when working with tabular data.

Since the weight decay plays an important part in the super-convergence phenomena, we do an experiment to find the optimal value and to see how it influences the performance. By the suggestion of (Smith, 2018) we try weight decay values for 10^{-3} , 10^{-5} and 0. First we do a learning rate range test with the different weight decays to identify the best value. The results are displayed in Figure 4.9. From the range test it seems that a weight decay of 10^{-5} gives the best performance. To check that this result holds during a full training run, we do an experiment to compare the full training runs with the different weight decays. The results are displayed in Figure 4.10. From the results it shows that the learning rate range test is a good indicator of the optimal weight decay, since in the full training run, the model with weight decay of 10^{-5} performed




Figure 4.8: Constant learning rate vs the 1Cycle schedule.

the best. Note, that the differences between the training runs are small, and thus if we were to choose one of the other weight decays for this task, we would have only performed slightly worse.




Figure 4.9: A learning rate range test with different weight decays.




Figure 4.10: A full training run with different weight decays.

Chapter 5

Experiments

"For us, the most important part of rigor is better empiricism, not more mathematical theories."

— Ali Rahimi and Ben Recht, *NIPS 2017*

5.1 Introduction

Since theory and practice does not always go hand-in-hand, it is usually advantageous to compliment a theoretical study or literature review with empirical results. Another motivation for empirical study is that we regard the ability to implement an approach equally as important as understanding the theory behind it. We characterise a good empirical experiment as one that is *rigorous* and *reproducible*. Recently the field of DL has been criticised for the growing gap between the understanding of its techniques and its practical successes¹ where most of the recent focus was on the latter. The speakers urged the deep learning community to be more rigorous in their experiments where, for them, the most important part of rigor is better empiricism, not more mathematical theories.

In this chapter we aim for good empiricism by exploring many hyperparameters in various data scenarios and doing cross-validation for unbiased performance measures along with standard errors. Our work is not necessarily about beating the benchmark and instead consists of simple experiments that aid in the understanding of how the techniques work and what effect they have.

Furthermore, we want all our experiments to be as reproducible as possible. Therefore we provide all the code, data and necessary documentation

¹How do I cite the talk given at NIPS2017 - <https://www.youtube.com/watch?v=Qi1Yry33TQE>

to reproduce the experiments that were done in this thesis (see Appendix C). This is often an overlooked feature of experiments, but is however crucial for transparent and accountable reporting and making your work useful for others to build upon.

The main aim of this chapter is to better understand the behaviours of certain models and parameters and to cross-check the literature with empirical observations. We focus on the same main modelling challenges for deep learning on tabular datasets that we discussed in chapter 4, which are:

- how to represent the inputs,
- how to learn from feature interactions, and
- how to be more sample efficient.

The model interpretation challenge was covered in section 4.5.

The more general hyperparameters, like learning rate, batch size, layer size and layer depth is not the focus of these experiments and some we have already experimented with in previous chapters. However, since these parameters are tightly linked with each other and other model parameters, we still do a hyperparamter search where we deem appropriate and report the findings in Appendix B. The rest of the chapter continues as follows: In §5.2 we discuss the dataset used for these experiments and why it was chosen. Thereafter, in §5.3, we discuss the general methodology of the experiments and how they are evaluated. Then the main experiments follow. §5.4 compares the performance of various entity embedding sizes. §5.5 looks at the proposed approaches to model high-order feature interactions. §5.6 evaluates the different approaches to avoid overfitting in constrained data environments.

5.2 Datasets

We chose to do our experiments on the Adult dataset for the following reasons:

- **Simplicity:** We wanted a dataset that is representative of a real-world case but one that does not have specific modelling challenges like plenty of missing values or highly imbalanced classes. Our goal is to evaluate the models on a generic tabular dataset and not one that requires special attention or the skills of a domain expert. Thus when we fall short we know it is because of the model and not something that is in the data.

- **Minimal preprocessing:** We want to focus our time on training the algorithms and not preprocessing the data. The Adult dataset is relatively clean.
- **Open access:** Since we want our work to be reproducible, we want the dataset we use to be accessible by all.
- **Good size:** Neural networks are data hungry and therefore for optimal performance we prefer a medium size dataset. If we want to test the performance of the models on smaller datasets, we can just run the experiment on a subset of the data.
- **Strong baselines:** In order to know how well we are doing we need to be able to compare our performance with those of others.

For more specifics on the dataset, one can refer to Appendix A

5.3 General Methodology

5.3.1 Loss Function and Evaluation Metric

Since this is a binary classification task we train the neural network to optimise the binary cross-entropy loss as defined in §2.3.3, with $K = 2$. This is the standard loss function to optimise for binary classification (Song *et al.*, 2018, Wang *et al.* (2017), Zhang *et al.* (2016), Qu *et al.* (2016)). We also monitor the accuracy of the model during training since this is often more interpretable, but note that the accuracy is not directly being optimised and that the binary cross-entropy only acts as an differentialable proxy of the accuracy.

5.3.2 Cross-validation

For most of the experiments we will do a 5-fold cross validation (Hastie *et al.*, 2009, p. 241) to estimate the performance of a model. That is, randomly dividing the dataset in five equal parts and then in turn, hold out one of those parts for validation purposes and train the model using the remaining four parts. Figure 5.1 visually explains how the dataset is sub-divided. The performance of the model can then be evaluated on the held-out part. This process is repeated for every one of the five segmentations of the dataset and thus five measurements of the performance of model is obtained. We can then compute the average over these five measurements to obtain a less biased estimate of

the model performance. Another advantage of this approach is that we can obtain standard error for the model performance. These standard errors will be displayed on the figures illustrating the results of the experiments as confidence intervals, *i.e.* $\mu \pm \sigma$. This helps us to evaluate the significance of the results.




Figure 5.1: 5-Fold Cross-validation dataset split schematic.

It is common in deep learning research to do a once-off split of the data dividing it into training, validation and (sometimes) testing datasets, for example in (Klambauer *et al.*, 2017, Song *et al.* (2018), Zhang *et al.* (2016)). Cross-validation is rarely done in Deep Learning, since the models typically take very long to train and any repetition is thus more costly. The problem with doing a once-off split is that it does not account for the variance of the model and the performance of the model can in fact be very sensitive to the subset of data. By doing cross-validation, we can have more robust performance metrics, including the benefit of reporting on standard errors. However, Deep Learning is also mostly applied to large datasets and if a large test set is available, the gains from cross-validation diminishes. Fortunately, the neural networks applied to tabular data are much smaller than ones used for unstructured data and therefore cross-validation makes sense.

5.3.3 Preprocessing

For our experiments we aimed to do as little feature engineering and preprocessing as possible, in the spirit of automated analysis. We do not create any new feature combinations. The only necessary preprocessing steps are to convert the categorical features into integers (necessary for entity embeddings), standard normal scaling for the continuous features (necessary for optimisation algorithms), mean impute the missing values in the continuous features and assign a “null” category to missing values in the categorical features. We do not do feature selection since we would want the model to learn by itself which features are relevant.

5.3.4 General Hyperparameters

Based on our findings in §4.6 we decide to train all models using the 1Cycle policy and the Adam optimiser. Unfortunately, we cannot follow the hyperparameter selection process suggested by (Smith, 2018) for all the experiments. The process is too manual. Therefore we follow the approach once to find a good selection of the learning rate, number of epochs and weight decay and then use these parameters for the rest of the experiments on this dataset. Some of this is reported in Appendix B. If the model significantly changes over experiments or we experience unstable training, we rerun the parameter selection process. Thus we will definitely not find the optimal model for each experiment but it should be sufficient to use for comparison purposes. According to (Smith, 2018) these parameters are also quite robust and the model is not too sensitive on these choices.

We follow a bit of a greedy approach when selecting optimal parameters. All the hyperparameters are very dependent on each other but we cannot run experiments for every possible combination. Therefore we find optimal parameters for a certain experiment and then assume that these parameters are also good for other experiments. If not specified otherwise, we used a simple MLP as the network mapping the input representation to the output. The basic structural hyperparameters we tested can be found in Appendix B. There we experimented with things like the dropout ratio, width and depth of the network, amongst other.

5.4 Input Representation

5.4.1 Embeddings Sizes

The aim of these experiments is to evaluate performance at different embedding sizes. We first explore fixed embedding sizes for all categorical features. The sizes we experiment with are [2,4,8]. Again, we expect there to be optimal embedding size for each variable depending on the cardinality of the variable and how complex its relationship is with the other variables and the target. We expect the ideal embedding size to be as small as possible but still being able to capture all of the information of the variable. But the only way to tell is by trial and error. We would also like to see if the rules-of-thumb mentioned in §4.2.2 provide reasonable performance. A benefit that these schemes can add is that they depend on the cardinality of the categorical feature.

The results of the one sized embeddings are displayed in Figure 5.2. We see that if all embeddings are fixed at size 2, the model showed greater performance than at 4 and 8. One reason for this may be that the categorical features are not that important for the prediction of the target and that a smaller representation helps to reduce noise.




Figure 5.2: Effect of the embedding size if all categorical features are mapped to the same number of dimensions.

We now compare the constant size 2 embeddings with the rules from (Wang *et al.*, 2017) and (de Brébisson *et al.*, 2015) to see how the cardinality dependency effects the performance. The results are shown in Figure 5.3. There is almost no difference between the approaches. Brebisson's methos performs

sightly better on the validation loss and the fixed size of 2 performs the best on accuracy. On other datasets this result will most likely be different if the categorical features contain more rich information.




Figure 5.3: Effect of variable sizes on the performance of the model.

5.5 Feature Interactions

5.5.1 Attention

We believe attention is the most promising approach to modelling high-order interactions between features for tabular data, since it has been so effective in other data domains. To test this assumption we implemented a multi-head attention module as described in (Song *et al.*, 2018) so that we can compare it to a standard neural network. We did not have the computing power to try different configurations for this module and thus we went with the biggest configuration that could be run in reasonable time. We chose $H = 3$ (Number of heads) and use embedding sizes of 3 all-round. We only used one multi-head attention layer and connected its output to a single hidden layer of size 200, which in turn is connected to the output layer. We compared it to a simple MLP with a similar number of total parameters. The results of this experiment is given in Figure 5.4. The model with the attention mechanism performs worse than the one without and exhibits higher variance. In the future we would like to experiment with more of the hyperparameters of the multi-head attention

module and try stacked attention layers. It will be interesting to see how this results would change if the dataset had more features.




Figure 5.4: Comparing the attention mechanism with a simple MLP.

5.5.2 SeLU

Deeper networks can help us learn higher-order feature interactions. The SeLU activation function is supposed to help us train deeper neural networks. Therefore we compare it to the ReLU activation function over two models, one with 2 layers and the other with 8, to investigate its impact on performance. The SeLU activation function was implemented with its special weight initialisation, but no dropout was used. In Figure 5.5 and Figure 5.6 we observe that the model with 8 layers and SeLU activations, perform the best overall in terms of the validation loss, but in terms of accuracy the ReLU model with 2 layers showed the best results. The differences observed are minor which leads us to believe that the choice is not that important and that sticking with shallow ReLU networks are sufficient until more experiments are done.

5.5.3 Skip-Connections

Another mechanism we can use to make deeper networks possible is a skip-connection. Skip-connections may also help to combine different orders of feature interactions. The skip connection combines the activations before and after a linear layer through an elementwise addition. We tested networks with skip-connections for a shallow and a deep neural network. The results




Figure 5.5: The average performance of ReLU and SeLU activation functions for shallow and deep networks over every epoch.




Figure 5.6: The average performance of ReLU and SeLU activation functions for shallow and deep networks.

are summarised in Figure 5.7 and Figure 5.8. On average did the deeper neural network perform better and in terms of the validation loss did the skip-connections slightly improve the performance. Again, there is very little between the two types and thus we cannot say one is better than the other. For simplicity sake, we would suggest not using skip-connections, until more experiments have been done.

5.6 Sample Efficiency

We did an experiment to see how the number of samples available to the network for training influences its learning ability. We trained the network at training set sizes of 1000, 2000, 4000, 8000, 16000 and 32000. The results are reported in Figure 5.9. As expected, the network performs better as the training set




Figure 5.7: Average performance at each epoch for shallow and deep neural networks, with and without skip connections.




Figure 5.8: Overall performance of the skip-connections used in a shallow and deep neural network.

increases, but the increase diminish as the sample size gets higher. We do observe a strange results where the model trained only using 4000 samples exhibit one of the best accuracies.

5.6.1 Data Augmentation

Mixup was the most promising form of data augmentation from our literature review. In this experiment we want to investigate how mixup with different mixup ratios influence the performance of our neural networks on the adult dataset. Since data augmentation also acts as a form of regularisation, we want to see the interaction between weight decay and mixup ratios. We experimented with mixup α -ratios of $[0,0.4]$ and with weight decays of $[10^{-5}, 10^{-3}]$ on a neural network with 3 hidden layers of 200 units each. The results of the experiment are captured in Figure 5.10 and Figure 5.11.

The results indicate that mixup is not improving the validation loss or the accuracy of the models. It is also interesting to see the interaction between




Figure 5.9: Effect of the number of training samples on the performance of neural networks.

mixup and weight decay. Since both are forms of regularisation we would expect one to reduce the need of the other, but this is not what the results suggest. Although a promising technique from the literature, we think the value of mixup depends on the dataset and whether or not interpolating between samples makes sense.




Figure 5.10: Average performance of models with various mixup and weight decay parameters.

5.6.2 Unsupervised Pretraining

Here we investigate whether or not DAEs as a method of unsupervised pre-training is beneficial to a supervised learning neural network. We trained a DAE with swap noise for 15 epochs on the Adult dataset, with the swap noise proportion set at 15%. The DAE had 3 hidden layers of 500 units each.




Figure 5.11: Performance per epoch for models with different weight decays and mixup ratios.

Then we transfer the learned weights to a supervised learning network, also with 3 hidden layers of 500 units. But the supervised network has a different output layer to the DAE. Therefore we first keep the transferred weights fixed and train only the last layer of the supervised learning network, as a way of initialising its weights. Otherwise if we would have trained all the weights together from the start, the random weights of the output layer might have interfered with the learned weights in the hidden layers. Once the output layer is initialised with this process, we can train all of the network simultaneously. This final stage of training is showed in Figure 5.12 where we compare it to a model without pretraining. We see that in terms of both the accuracy and the validation loss that the pretrained model has an advantage over the trained from scratch classifier. This makes sense since the pretrained model does not have to start from a random initialisation. But as training continues this gaps becomes smaller. In terms of the validation loss the two models are equivalent when training completed, but in terms of the accuracy the trained from scratch classifier outperformed the pretrained model. We do not believe these results are conclusive since there are still so many avenues to explore for unsupervised pretraining with DAEs. We do not know yet how to design the DAE in terms of its network architecture, amount of noise injected, lenght of training cycle, etc, and then what the best way is of transferring this knowledge to a classifier. As a test to validate if the DAE learned something useful one can feed the activations from one of the DAE hidden layers to another machine learning algorithm and see whether these representations perform better than with the

raw data. From the results of this basic implementation, we do still believe that there is value in this approach and that it warrants further exploration.




Figure 5.12: The effect of pretraining on the classifier’s performance.

5.7 General Thoughts

We ran thorough experiments for the most promising techniques in deep learning for tabular data. In most of the experiments there were not clear differences between the performances of the different techniques. Our one hypothesis is that the Adult dataset does not lend itself to more advanced modelling techniques and that the basic neural network is close to the best one can get. Our other concern is that the greedy fashion of our hyperparameter selection has limited the various approaches. Since it was not feasible to tune the required parameters for all models in all folds, we selected a set of parameters from initial experiments and used them for most of the experiments. But since these parameters are so inter-connected, changing one already means that the others are not optimal anymore. We were aware of this limitation from the start, but we still do not know however to what extent it influenced our results.

Chapter 6

Conclusion

In this work we set out to investigate a relatively under-explored area of deep learning; applying deep learning approaches to tabular datasets. The goal was to find and understand the best approaches for this task by reviewing the literature and doing an empirical study. To assist in this endeavour we also aimed to review the core concepts of neural networks and discuss the modern advancements in deep learning, with the hope of identifying approaches that are transferrable to our domain.

In this chapter we summarise the work of this thesis and evaluate how each of the chapters contributed to the set out objectives (§6.1). The limitations of this work is discussed in §6.2, how they affected the study and how it could be improved upon. We conclude this work with a section on promising future directions in this area of research (§6.3).

6.1 Summary

In **Chapter 1** the motivation and the objectives of this work were introduced. It was stated that deep learning for tabular data is an important topic but that it has not received enough attention in the literature. The main issues that (we feel) need to be addressed in order to make progress in the field were highlighted. The chapter also provided an overview of the fundamentals of Statistical Learning Theory, including different learning paradigms, loss functions, optimisation and overfitting, to serve as background for the problems we were trying to solve.

Chapter 2 was about neural networks and how to train them. Neurons,

layers, activation functions and network architectures were described in detail to gain a complete understanding of the mechanics of neural networks. To describe how a neural network is trained, the backpropagation and stochastic gradient descent algorithms were introduced, with the help of examples. The chapter included a brief look at basic regularisation methods for neural networks. This was followed by a section on representation learning to gain insight to what a neural network is actually learning and how it does so. This chapter provided us with the fundamentals of neural networks to enhance our capacity to understand the more advanced developments in deep learning that follows.

Chapter 3 was about deep learning in general. The major developments contributing to the recent success of deep learning were investigated. The chapter introduced autoencoders and how they can be used for unsupervised pre-training, along with the concept of transfer learning. Very effective regularisation techniques, data augmentation and dropout were described. This was followed by a review of the more modern layers and architecture designs in deep learning, which included normalisation layers, skip-connections, embeddings and the attention mechanism. The chapter included a section on the concept of superconvergence and how it can be achieved by using the 1Cycle policy and more effective hyperparameter selection, with the help of an example and experiment. A brief discussion on the neural network specific and model agnostic tools available for interpreting deep neural networks was also provided and examples were shown. The topics that were focussed on in this chapter gave insight into how deep learning achieved its success in other data domains, which provided a good starting point in an attempt to improve deep learning for tabular data.

Chapter 4 was about deep learning for tabular data. In this chapter a review was done on all the latest literature on applying deep neural networks to tabular data. The main challenges to deep learning for tabular data, as stated in Chapter 1, guided this review. The chapter explores the various ways of preprocessing continuous features and embedding categorical features in a way that is optimal for a neural network. Then it investigated how a neural network can be encouraged to learn better feature interactions, by using attention modules and cross-features. The SeLU activation was proposed, along with skip-connections, as ways to build deeper neural networks for more complex feature learning. This was followed by evaluating approaches that make deep neural networks more sample efficient. It focussed on denoising autoencoders

as a way of pretraining a network and the data augmentation techniques, swap noise and mixup. Reference was made to per-feature regularisation. Then there was a discussion on more ways of interpreting neural networks for tabular data, considering the approaches discussed earlier in the chapter. An example of interpreting a network with knowledge distillation was also given. The chapter concluded with an empirical investigation of superconvergence on tabular data. This chapter provided us with a solid technical understanding of the approaches available to overcome the challenges in the field. It also gave sufficient detail to be able to implement these promising approaches.

Chapter 5 contains the empirical experiments. The experiments served to assist the exploratory study of deep learning for tabular data. The experiments were categorised by three main themes: input representation, feature interaction and sample efficiency. In the input representation experiments we evaluated the effect the entity embedding sizes had on the performance of the model. The attention mechanism, along with skip-connections and the SeLU activation were experimented with to see how they do in modeling feature interactions. Finally, in the sample efficiency section, we investigated how sensitive neural networks are to the number of observations in a dataset and tested unsupervised pretraining as an approach to alleviate this sensitivity. We also looked at the mixup and swap noise data augmentation approaches as means to avoid overfitting.

6.2 Limitations

There were various aspects that limited the impact of this study. Two technical limitations were experienced:

- **Access to large compute:** Deep learning techniques are notorious for the amount of computing power they require. There was limited access to cloud computing providers on which some of the experiments were done. However, in most cases, we only had access to a small personal machine without a graphical processing unit (GPU). This significantly increased the running time of the experiments and hindered rapid iteration of different approaches. In future work we would make sure that sufficient computing power is available for experimenting in this field.

- **Access to quality code:** Many of the recent developments discussed in this work had no official implementation at the time of writing. This forced us to rewrite many of the code in order to validate the results claimed in the literature and to be able to use it in our data and models. Sometimes technical details were left out of the original papers and we had to improvise to get a working example. By doing it this way we risked using code for our experiments that might not have been exactly what the authors intend.

The other limitations to this study were:

- **Experiments on a single dataset:** Due to the technical limitations described above, we only had the capacity to process a single dataset for our experiments. This made our results inconclusive, but still adds value as an exploratory study. There is also a possibility that this dataset was in a sense too easy to solve and that we could not demonstrate the full power of the deep learning approaches. This study would have benefitted from experiments on datasets with different properties and with different tasks. We published all of our code online if anyone wants to build on this work.
- **Based on pre-prints:** Deep learning is such a fast developing area of research and in an attempt to keep this work relevant, pre-prints of publications were cited. Pre-prints are not peer-reviewed and subject to change. We did our best to critically evaluate the work cited and to confirm findings with experiments. Although we tried to keep up with the deluge of publications, there remains a possibility that new publications arised during the post-research phase of this work.

6.3 Promising Future Directions

Throughout this thesis we have identified promising research directions which, however, did not fit into the scope of this work. Besides repeating the empirical work on a wider variety of datasets, we propose the following to be a valuable research path to follow:

Generative models for tabular data. Here we specifically refer to *Variational Autoencoders* (VAE) (Kingma and Welling, 2013) and *Generative*

Adversarial Networks (GAN) (Goodfellow *et al.*, 2014). A VAE provides a probabilistic manner for describing an observation in a latent space. That is, instead of using a single value to describe an attribute in the latent space, like the standard autoencoder, a VAE uses a probability distribution as a description of a latent space attribute. The greatest value we think VAEs can add is as a method of unsupervised pretraining. When using DAEs for unsupervised pretraining, it needs to be injected with some noise. However, the best noise schemes we found, were swap noise and blank-out, but as mentioned before, neither makes complete sense in the tabular dataset environment. A VAE can provide a more robust way of doing unsupervised learning because its decoding function needs to learn to process probabilistic output and thus it is not reliant on noise injection. In addition, once trained, a VAE can be used to generate new samples by sampling from the latent probabilistic distribution, which may serve as a way to generate more training samples.

We also believe that GANs offer a good alternative to generating new training samples. A GAN consists of two neural networks, a generator (G) and a discriminator (D). The task of G is: given random noise as input, generate artificial samples of the data that are indistinguishable from the genuine samples. The task D is to attempt to discriminate between the two. GANs have shown tremendous value in data synthesis, especially in the domains of computer vision and speech synthesis, producing life-like faces (Karras *et al.*, 2017) and voices (Donahue *et al.*, 2018). We think GANs can achieve similar success in data synthesis for tabular data and can thus also be used to artificially enlarge training datasets for supervised learning.

Appendices

Appendix A

Datasets

Details of each of the datasets used in Chapter 5 and elsewhere.

A.1 Adult Dataset

The Adult dataset, originally used in (Kohavi, 1996), is data that was extracted from the census bureau database¹ and can be accessed from this link². The task here is to predict whether or not a certain person's income exceeds \$50,000 per year. Thus it is posed as a binary classification problem.

In total there are 14 features and 48,842 observations. Two-thirds of the observations were randomly selected to form the training set and the rest allocated to the test set. Note, that we will not use any of the observations in the test set during our experiments. Hyperparameter and model selection decisions are made based on the validation dataset performance and then if we wish we can evaluate the selected models on the test set for the most accurate estimation of the generalisation ability of the models.

The details for each of the features are listed below. We indicate the continuous features and the classes for the categorical features.

- *age*: continuous
- *workclass*: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked
- *fnlwgt*: continuous

¹<http://www.census.gov/en.html>

²<http://archive.ics.uci.edu/ml/datasets/Adult>

- *education*: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool
- *education-num*: continuous
- *marital-status*: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse
- *occupation*: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces
- *relationship*: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried
- *race*: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black
- *sex*: Female, Male
- *capital-gain*: continuous
- *capital-loss*: continuous
- *hours-per-week*: continuous
- *native-country*: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinadad&Tobago, Peru, Hong, Holand-Netherlands

The distributions of the continuous features are plotted in Figure A.1 and the counts of the categories for each of the categorical features can be found in Figure A.2. The highest reported accuracy in the original paper (Kohavi, 1996) was 85.9% by the NBTree algorithm and in an unpublished experiment³ one researcher achieve ~88% accuracy with boosted trees.

³<https://www.kaggle.com/kanav0183/catboost-and-other-class-algos-with-88-accuracy>




Figure A.1: Histograms for each of the continuous features in the Adult dataset.




Figure A.2: Bar plot for each of the categorical features in the Adult dataset.

Appendix B

Hyperparameter Search

B.1 Width and Depth of Network

This is a very common hyperparameter to train, for example, done by (Guo *et al.*, 2017, Qu *et al.* (2016), Zhang *et al.* (2016)). Here we investigate the effect of the size of the network on the different datasets. We compare the performance of the models at different numbers and sizes of layers. Larger networks are more flexible and therefore we expect it to act similarly to any learning model flexibility parameter. Increasing the network size will be beneficial up until a certain point until it becomes too big and be more prone to overfitting. We also want to get a feel for how important these hyperparameters are. On the Adult dataset we experimented with layer depths of 1, 3, 6 and 12, and layer widths of 32, 128, 512, 2048. Note give all layers the same width as this is found to work good enough in practice (Guo *et al.*, 2017, Qu *et al.* (2016), Zhang *et al.* (2016)). The results of the experiment are displayed in Figure B.1. The results are mixed between the accuracy and the validation loss and there is not much separating any of the models. Except for the very wide network that had a diverging validation loss. For the sake of simplicity we would recommend to use a 3-layer network with 128 units.

B.2 Dropout

Dropout is almost always used in deep learning and it is a typical parameter to tune before modelling. We wanted to find out what the best dropout ratio is for our models and how the performance varies over different values. We




Figure B.1: Effect of the layer width and network depth on the performance on the Adult dataset.

tested dropout proportions of $[0.1, 0.2, \dots, 0.8]$ on a wide and narrow network respectively. The results are given in Figure B.2. From these results we see that any dropout ratio between 0.2 and 0.5 will suffice for both types of networks.




Figure B.2: The effect of dropout on wide and narrow neural networks.

Appendix C

Software and Code

C.1 Development Environment

- Deep Learning Libraries: Pytorch and Fastai
- Hardware:
 - MacBook Air (2017)
 - Cloud Providers: AWS EC2¹, SalamanderAI², Google Compute Engine³.
- Programming Language: Python
- github for version control
- RMarkdown for writing and compiling the thesis document

C.2 Code and Reproducibility

Note that all of the code used in the thesis, including the source documents, is made available in the tabularLearner Github repository ⁴. More instructions on how to implement the code is contained in the file named `README.md`, in the repository.

¹<https://aws.amazon.com/ec2/>

²<https://salamander.ai/>

³<https://cloud.google.com/compute/>

⁴<https://github.com/jandremarais/tabularLearner>

Bibliography

- Alain, G. and Bengio, Y. (2014). What regularized auto-encoders learn from the data-generating distribution. *The Journal of Machine Learning Research*, vol. 15, no. 1, pp. 3563–3593.
- Anonymous (2019). Tabnn: A universal neural network solution for tabular data. In: *Submitted to International Conference on Learning Representations*. Under review.
Available at: <https://openreview.net/forum?id=r1eJssCqY7>
- Ba, L.J. and Caurana, R. (2013). Do deep nets really need to be deep? *CoRR*, vol. abs/1312.6184. 1312.6184.
Available at: <http://arxiv.org/abs/1312.6184>
- Bahdanau, D., Cho, K. and Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. *CoRR*, vol. abs/1409.0473. 1409.0473.
Available at: <http://arxiv.org/abs/1409.0473>
- Battenberg, E., Chen, J., Child, R., Coates, A., Gaur, Y., Li, Y., Liu, H., Satheesh, S., Seetapun, D., Sriram, A. and Zhu, Z. (2017). Exploring neural transducers for end-to-end speech recognition. *CoRR*, vol. abs/1707.07413. 1707.07413.
Available at: <http://arxiv.org/abs/1707.07413>
- Bengio, Y., Boulanger-Lewandowski, N. and Pascanu, R. (2012). Advances in optimizing recurrent networks. *CoRR*, vol. abs/1212.0901. 1212.0901.
Available at: <http://arxiv.org/abs/1212.0901>
- Bengio, Y., Courville, A. and Vincent, P. (2013 Aug). Representation learning: A review and new perspectives. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 8, pp. 1798–1828. ISSN 0162-8828.
- Breiman, L. (2001). Random forests. *Machine learning*, vol. 45, no. 1, pp. 5–32.
- Chapelle, O., Weston, J., Bottou, L. and Vapnik, V. (2001). Vicinal risk minimization. In: Leen, T.K., Dietterich, T.G. and Tresp, V. (eds.), *Advances in Neural*

- Information Processing Systems 13*, pp. 416–422. MIT Press.
Available at: <http://papers.nips.cc/paper/1876-vicinal-risk-minimization.pdf>
- Che, Z., Purushotham, S., Khemani, R.G. and Liu, Y. (2016). Interpretable deep models for icu outcome prediction. *AMIA ... Annual Symposium proceedings. AMIA Symposium*, vol. 2016, pp. 371–380.
- Cheng, H., Koc, L., Harmsen, J., Shaked, T., Chandra, T., Aradhye, H., Anderson, G., Corrado, G., Chai, W., Ispir, M., Anil, R., Haque, Z., Hong, L., Jain, V., Liu, X. and Shah, H. (2016a). Wide & deep learning for recommender systems. *CoRR*, vol. abs/1606.07792. 1606.07792.
Available at: <http://arxiv.org/abs/1606.07792>
- Cheng, J., Dong, L. and Lapata, M. (2016b). Long short-term memory-networks for machine reading. *CoRR*, vol. abs/1601.06733. 1601.06733.
Available at: <http://arxiv.org/abs/1601.06733>
- Choromanska, A., Henaff, M., Mathieu, M., Arous, G.B. and LeCun, Y. (2014). The loss surface of multilayer networks. *CoRR*, vol. abs/1412.0233. 1412.0233.
Available at: <http://arxiv.org/abs/1412.0233>
- Clevert, D.-A., Unterthiner, T. and Hochreiter, S. (2015 November). Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs). *ArXiv e-prints*. 1511.07289.
- Covington, P., Adams, J. and Sargin, E. (2016). Deep neural networks for youtube recommendations. In: *Proceedings of the 10th ACM Conference on Recommender Systems*. New York, NY, USA.
- Cybenko, G. (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of control, signals and systems*, vol. 2, no. 4, pp. 303–314.
- de Brébisson, A., Simon, É., Auvolat, A., Vincent, P. and Bengio, Y. (2015). Artificial neural networks applied to taxi destination prediction. *CoRR*, vol. abs/1508.00021. 1508.00021.
Available at: <http://arxiv.org/abs/1508.00021>
- Devlin, J., Chang, M.-W., Lee, K. and Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- Donahue, C., McAuley, J. and Puckette, M. (2018). Synthesizing audio with generative adversarial networks. *arXiv preprint arXiv:1802.04208*.

- Duchi, J., Hazan, E. and Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research*, vol. 12, no. Jul, pp. 2121–2159.
- Duong, L., Anastasopoulos, A., Chiang, D., Bird, S. and Cohn, T. (2016). An attentional model for speech translation without transcription. In: *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 949–959.
- Erhan, D., Bengio, Y., Courville, A. and Vincent, P. (2009). Visualizing higher-layer features of a deep network.
- Fernández-Delgado, M., Cernadas, E., Barro, S. and Amorim, D. (2014). Do we need hundreds of classifiers to solve real world classification problems? *Journal of Machine Learning Research*, vol. 15, pp. 3133–3181.
Available at: <http://jmlr.org/papers/v15/delgado14a.html>
- Fridman, L., Brown, D.E., Glazer, M., Angell, W., Dodd, S., Jenik, B., Terwilliger, J., Kindelsberger, J., Ding, L., Seaman, S., Abraham, H., Mehler, A., Sipperley, A., Pettinato, A., Seppelt, B., Angell, L., Mehler, B. and Reimer, B. (2017). MIT autonomous vehicle technology study: Large-scale deep learning based analysis of driver behavior and interaction with automation. *CoRR*, vol. abs/1711.06976. 1711.06976.
Available at: <http://arxiv.org/abs/1711.06976>
- Friedman, J.H. (2001). Greedy function approximation: a gradient boosting machine. *Annals of statistics*, pp. 1189–1232.
- Friedman, J.H. and Stuetzle, W. (1981). Projection pursuit regression. *Journal of the American statistical Association*, vol. 76, no. 376, pp. 817–823.
- Frosst, N. and Hinton, G.E. (2017). Distilling a neural network into a soft decision tree. *CoRR*, vol. abs/1711.09784. 1711.09784.
Available at: <http://arxiv.org/abs/1711.09784>
- Gatys, L.A., Ecker, A.S. and Bethge, M. (2015). A neural algorithm of artistic style. *CoRR*, vol. abs/1508.06576. 1508.06576.
Available at: <http://arxiv.org/abs/1508.06576>
- Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*. MIT Press. <http://www.deeplearningbook.org>.

- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A. and Bengio, Y. (2014). Generative adversarial nets. In: *Advances in neural information processing systems*, pp. 2672–2680.
- Guo, C. and Berkhahn, F. (2016). Entity embeddings of categorical variables. *CoRR*, vol. abs/1604.06737. 1604.06737.
Available at: <http://arxiv.org/abs/1604.06737>
- Guo, H., Tang, R., Ye, Y., Li, Z. and He, X. (2017). Deepfm: A factorization-machine based neural network for CTR prediction. *CoRR*, vol. abs/1703.04247. 1703.04247.
Available at: <http://arxiv.org/abs/1703.04247>
- Haldar, M., Abdool, M., Ramanathan, P., Xu, T., Yang, S., Duan, H., Zhang, Q., Barrow-Williams, N., Turnbull, B.C., Collins, B.M. and Legrand, T. (2018 October). Applying Deep Learning To Airbnb Search. *ArXiv e-prints*. 1810.09591.
- Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The elements of statistical learning: data mining, inference and prediction*. 2nd edn. Springer.
Available at: <http://www-stat.stanford.edu/~tibs/ElemStatLearn/>
- He, K., Zhang, X., Ren, S. and Sun, J. (2015a). Deep residual learning for image recognition. *CoRR*, vol. abs/1512.03385. 1512.03385.
Available at: <http://arxiv.org/abs/1512.03385>
- He, K., Zhang, X., Ren, S. and Sun, J. (2015b). Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *CoRR*, vol. abs/1502.01852. 1502.01852.
Available at: <http://arxiv.org/abs/1502.01852>
- Hinton, G., Vinyals, O. and Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*.
- Hinton, G.E., Osindero, S. and Teh, Y.-W. (2006 July). A fast learning algorithm for deep belief nets. *Neural Comput.*, vol. 18, no. 7, pp. 1527–1554. ISSN 0899-7667.
Available at: <http://dx.doi.org/10.1162/neco.2006.18.7.1527>
- Hinton, G.E. and Salakhutdinov, R.R. (2006). Reducing the dimensionality of data with neural networks. *science*, vol. 313, no. 5786, pp. 504–507.
- Hinton, G.E., Srivastava, N., Krizhevsky, A., Sutskever, I. and Salakhutdinov, R. (2012). Improving neural networks by preventing co-adaptation of feature detectors. *CoRR*, vol. abs/1207.0580.
Available at: <http://arxiv.org/abs/1207.0580>

- Howard, J. and Ruder, S. (2018). Fine-tuned language models for text classification. *arXiv preprint arXiv:1801.06146*.
- Hu, J., Shen, L. and Sun, G. (2017). Squeeze-and-excitation networks. *CoRR*, vol. abs/1709.01507. 1709.01507.
Available at: <http://arxiv.org/abs/1709.01507>
- Huang, F., Ash, J., Langford, J. and Schapire, R. (2017). Learning deep resnet blocks sequentially using boosting theory. *arXiv preprint arXiv:1706.04964*.
- Huang, G., Liu, Z. and Weinberger, K.Q. (2016). Densely connected convolutional networks. *CoRR*, vol. abs/1608.06993. 1608.06993.
Available at: <http://arxiv.org/abs/1608.06993>
- Ioffe, S. and Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. *CoRR*, vol. abs/1502.03167.
Available at: <http://arxiv.org/abs/1502.03167>
- Karras, T., Aila, T., Laine, S. and Lehtinen, J. (2017). Progressive growing of gans for improved quality, stability, and variation. *CoRR*, vol. abs/1710.10196. 1710.10196.
Available at: <http://arxiv.org/abs/1710.10196>
- Kingma, D.P. and Ba, J. (2014). Adam: A method for stochastic optimization. *CoRR*, vol. abs/1412.6980. 1412.6980.
Available at: <http://arxiv.org/abs/1412.6980>
- Kingma, D.P. and Welling, M. (2013). Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*.
- Klambauer, G., Unterthiner, T., Mayr, A. and Hochreiter, S. (2017). Self-normalizing neural networks. *CoRR*, vol. abs/1706.02515. 1706.02515.
Available at: <http://arxiv.org/abs/1706.02515>
- Kohavi, R. (1996). Scaling up the accuracy of naive-bayes classifiers: a decision-tree hybrid. Citeseer.
- Kosar, R. and Scott, D.W. (2018 January). The Hybrid Bootstrap: A Drop-in Replacement for Dropout. *ArXiv e-prints*. 1801.07316.
- Krizhevsky, A., Sutskever, I. and Hinton, G.E. (2012). Imagenet classification with deep convolutional neural networks. In: *Proceedings of the 25th International Conference on Neural Information Processing Systems*, NIPS'12, pp. 1097–1105.

Curran Associates Inc., USA.

Available at: <http://dl.acm.org/citation.cfm?id=2999134.2999257>

Larochelle, H., Bengio, Y., Louradour, J. and Lamblin, P. (2009). Exploring strategies for training deep neural networks. *Journal of machine learning research*, vol. 10, no. Jan, pp. 1–40.

Lecun, Y., Bengio, Y. and Hinton, G. (2015 5). Deep learning. *Nature*, vol. 521, no. 7553, pp. 436–444. ISSN 0028-0836.

Li, M., Zhang, T., Chen, Y. and Smola, A.J. (2014). Efficient mini-batch training for stochastic optimization. In: *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '14, pp. 661–670. ACM, New York, NY, USA. ISBN 978-1-4503-2956-9.

Available at: <http://doi.acm.org/10.1145/2623330.2623612>

Lundberg, S. and Lee, S. (2017). A unified approach to interpreting model predictions. *CoRR*, vol. abs/1705.07874. 1705.07874.

Available at: <http://arxiv.org/abs/1705.07874>

Maas, A.L., Hannun, A.Y. and Ng, A.Y. (). Rectifier nonlinearities improve neural network acoustic models.

Mikolov, T., Sutskever, I., Chen, K., Corrado, G.S. and Dean, J. (2013). Distributed representations of words and phrases and their compositionality. In: *Advances in neural information processing systems*, pp. 3111–3119.

Miotto, R., Li, L., Kidd, B.A. and Dudley, J.T. (2016). Deep patient: An unsupervised representation to predict the future of patients from the electronic health records. In: *Scientific reports*.

Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D. and Riedmiller, M.A. (2013). Playing atari with deep reinforcement learning. *CoRR*, vol. abs/1312.5602. 1312.5602.

Available at: <http://arxiv.org/abs/1312.5602>

Mogren, O. (2016). C-RNN-GAN: continuous recurrent neural networks with adversarial training. *CoRR*, vol. abs/1611.09904. 1611.09904.

Available at: <http://arxiv.org/abs/1611.09904>

Perez, L. and Wang, J. (2017). The effectiveness of data augmentation in image classification using deep learning. *CoRR*, vol. abs/1712.04621. 1712.04621.

Available at: <http://arxiv.org/abs/1712.04621>

- Qu, Y., Cai, H., Ren, K., Zhang, W., Yu, Y., Wen, Y. and Wang, J. (2016). Product-based neural networks for user response prediction. *CoRR*, vol. abs/1611.00144. 1611.00144.
Available at: <http://arxiv.org/abs/1611.00144>
- Rajkomar, A., Oren, E., Chen, K., Dai, A.M., Hajaj, N., Liu, P.J., Liu, X., Sun, M., Sundberg, P., Yee, H., Zhang, K., Duggan, G.E., Flores, G., Hardt, M., Irvine, J., Le, Q.V., Litsch, K., Marcus, J., Mossin, A., Tansuwan, J., Wang, D., Wexler, J., Wilson, J., Ludwig, D., Volchenboum, S.L., Chou, K., Pearson, M., Madabushi, S., Shah, N.H., Butte, A.J., Howell, M., Cui, C., Corrado, G. and Dean, J. (2018). Scalable and accurate deep learning for electronic health records. *CoRR*, vol. abs/1801.07860. 1801.07860.
Available at: <http://arxiv.org/abs/1801.07860>
- Rosenblatt, F. (1962). *Principles of neurodynamics: perceptrons and the theory of brain mechanisms*. Report (Cornell Aeronautical Laboratory). Spartan Books.
Available at: <https://books.google.ca/books?id=7FhRAAAAMAAJ>
- Rumelhart, D.E., Hinton, G.E. and Williams, R.J. (1988). Neurocomputing: Foundations of research. chap. Learning Representations by Back-propagating Errors, pp. 696–699. MIT Press, Cambridge, MA, USA. ISBN 0-262-01097-6.
Available at: <http://dl.acm.org/citation.cfm?id=65669.104451>
- Sarikaya, R. (2017 Jan). The technology behind personal digital assistants: An overview of the system architecture and key components. *IEEE Signal Processing Magazine*, vol. 34, no. 1, pp. 67–81. ISSN 1053-5888.
- Selvaraju, R.R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., Batra, D. et al. (2017). Grad-cam: Visual explanations from deep networks via gradient-based localization. In: *ICCV*, pp. 618–626.
- Shavitt, I. and Segal, E. (2018 May). Regularization Learning Networks: Deep Learning for Tabular Datasets. *ArXiv e-prints*. 1805.06440.
- Shi, W., Caballero, J., Huszár, F., Totz, J., Aitken, A.P., Bishop, R., Rueckert, D. and Wang, Z. (2016). Real-time single image and video super-resolution using an efficient sub-pixel convolutional neural network. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 1874–1883.
- Shickel, B., Tighe, P., Bihorac, A. and Rashidi, P. (2017). Deep EHR: A survey of recent advances on deep learning techniques for electronic health record (EHR)

- analysis. *CoRR*, vol. abs/1706.03446. 1706.03446.
Available at: <http://arxiv.org/abs/1706.03446>
- Shimodaira, H. (2000). Improving predictive inference under covariate shift by weighting the log-likelihood function. *Journal of Statistical Planning and Inference*, vol. 90, no. 2, pp. 227 – 244. ISSN 0378-3758.
Available at: <http://www.sciencedirect.com/science/article/pii/S0378375800001154>
- Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., Hubert, T., Baker, L., Lai, M., Bolton, A. et al. (2017). Mastering the game of go without human knowledge. *Nature*, vol. 550, no. 7676, p. 354.
- Smith, L.N. (2015). No more pesky learning rate guessing games. *CoRR*, vol. abs/1506.01186. 1506.01186.
Available at: <http://arxiv.org/abs/1506.01186>
- Smith, L.N. (2018). A disciplined approach to neural network hyper-parameters: Part 1 - learning rate, batch size, momentum, and weight decay. *CoRR*, vol. abs/1803.09820. 1803.09820.
Available at: <http://arxiv.org/abs/1803.09820>
- Smith, L.N. and Topin, N. (2017). Super-convergence: Very fast training of residual networks using large learning rates. *CoRR*, vol. abs/1708.07120. 1708.07120.
Available at: <http://arxiv.org/abs/1708.07120>
- Song, W., Shi, C., Xiao, Z., Duan, Z., Xu, Y., Zhang, M. and Tang, J. (2018). AutoInt: Automatic feature interaction learning via self-attentive neural networks. *CoRR*, vol. abs/1810.11921. 1810.11921.
Available at: <http://arxiv.org/abs/1810.11921>
- Sun, Y., Wang, X. and Tang, X. (2014). Deeply learned face representations are sparse, selective, and robust. *CoRR*, vol. abs/1412.1265. 1412.1265.
Available at: <http://arxiv.org/abs/1412.1265>
- Van Der Maaten, L., Chen, M., Tyree, S. and Weinberger, K.Q. (2013). Learning with marginalized corrupted features. In: *Proceedings of the 30th International Conference on International Conference on Machine Learning - Volume 28*, ICML'13, pp. I–410–I–418. JMLR.org.
Available at: <http://dl.acm.org/citation.cfm?id=3042817.3042865>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L. and Polosukhin, I. (2017). Attention is all you need. *CoRR*, vol. abs/1706.03762.

1706.03762.

Available at: <http://arxiv.org/abs/1706.03762>

Vincent, P., Larochelle, H., Bengio, Y. and Manzagol, P.-A. (2008). Extracting and composing robust features with denoising autoencoders. In: *Proceedings of the 25th International Conference on Machine Learning*, ICML '08, pp. 1096–1103. ACM, New York, NY, USA. ISBN 978-1-60558-205-4.

Available at: <http://doi.acm.org/10.1145/1390156.1390294>

Wang, R., Fu, B., Fu, G. and Wang, M. (2017). Deep & cross network for ad click predictions. *CoRR*, vol. abs/1708.05123. 1708.05123.

Available at: <http://arxiv.org/abs/1708.05123>

Wu, Y., Schuster, M., Chen, Z., Le, Q.V., Norouzi, M., Macherey, W., Krikun, M., Cao, Y., Gao, Q., Macherey, K., Klingner, J., Shah, A., Johnson, M., Liu, X., Kaiser, L., Gouws, S., Kato, Y., Kudo, T., Kazawa, H., Stevens, K., Kurian, G., Patil, N., Wang, W., Young, C., Smith, J., Riesa, J., Rudnick, A., Vinyals, O., Corrado, G., Hughes, M. and Dean, J. (2016). Google’s neural machine translation system: Bridging the gap between human and machine translation. *CoRR*, vol. abs/1609.08144. 1609.08144.

Available at: <http://arxiv.org/abs/1609.08144>

Xu, K., Ba, J., Kiros, R., Cho, K., Courville, A.C., Salakhutdinov, R., Zemel, R.S. and Bengio, Y. (2015). Show, attend and tell: Neural image caption generation with visual attention. *CoRR*, vol. abs/1502.03044. 1502.03044.

Available at: <http://arxiv.org/abs/1502.03044>

Yosinski, J., Clune, J., Nguyen, A., Fuchs, T. and Lipson, H. (2015). Understanding neural networks through deep visualization. *arXiv preprint arXiv:1506.06579*.

Zeiler, M.D. and Fergus, R. (2014). Visualizing and understanding convolutional networks. In: *European conference on computer vision*, pp. 818–833. Springer.

Zhang, H., Cissé, M., Dauphin, Y.N. and Lopez-Paz, D. (2017). mixup: Beyond empirical risk minimization. *CoRR*, vol. abs/1710.09412. 1710.09412.

Available at: <http://arxiv.org/abs/1710.09412>

Zhang, W., Du, T. and Wang, J. (2016). Deep learning over multi-field categorical data: A case study on user response prediction. *CoRR*, vol. abs/1601.02376. 1601.02376.

Available at: <http://arxiv.org/abs/1601.02376>

- Zhou, B., Khosla, A., Lapedriza, A., Oliva, A. and Torralba, A. (2016). Learning deep features for discriminative localization. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 2921–2929.
- Zhou, G., Song, C., Zhu, X., Fan, Y., Zhu, H., Ma, X., Yan, Y., Jin, J., Li, H. and Gai, K. (2017 June). Deep Interest Network for Click-Through Rate Prediction. *ArXiv e-prints*. 1706.06978.